

PRACTICAL SESSION 2 Supervised DEEP LEARNING

The dataset is a *subset* of CIFAR10 (a very popular dataset in Machine Learning) with only 4 classes: **BIRD, CAT, FROG, HORSE**. We will call our dataset **CIFAR4**. It is composed of 24000 images from 4 types of animals.

Input data are square color images (RGB). The size of a picture is 32x32. So we have an input tensor of size [32,32,3]. Each pixel has a red, green and blue UINT8 values, i.e. in [0,255].



The dataset is available on [ECAMPUS](#)

The main objective is to train different deep neural network models to classify an image into one of the 4 classes.

The goal of this practical session is to experiment! Everything you need is easily available in the TENSORFLOW documentation that contains many examples.

BONUS STRIKE 🧠

You can make additional experiences. In that case, please add bonus works at the end of your notebook in the **[BONUS]** section.

Bonus works could bring bonus points on the global note for the practical sessions **if and only if** all the mandatory jobs have been done correctly.

✓ IMPORT LIBS

We will use TENSORFLOW, the Deep Learning platform from Google.

TensorFlow is easy to understand and the documentation and [tutorials](#) are (very) useful when learning Deep Learning.

I already know how to develop DL models and I want to use PyTorch instead. Is it possible?

==> YES. [PYTORCH](#) (from Meta) is the other big reference for DL libraries. Note that, in that case, *it's your choice* and you must be *self-sufficient* in case of development issues.

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import urllib
import zipfile
import matplotlib.pyplot as plt
from PIL import Image
import tensorflow as tf
```

✓ [WARNING] About the use of COLAB and GPU

On the top-right of the colab environment, you can chose the execution environment ("Modify the execution type") between CPU and GPU (Nvidia T4). With a free google account, there is no limitation with the CPU-only mode. With the GPU, you have a limitation that is dynamically set by Google regarding the overall load on their infrastructure. Usually, it could be 1 hour per day.

For this practical session, you can keep the CPU mode for the MLP part. For the CNN section, you can swith to the GPU mode if you consider the processing are too slow. But, be careful and not too GPU-enthusiast.

NB: You can check the CPU info with `!cat /proc/cpuinfo`

```
#Get CPU info
!cat /proc/cpuinfo
#Get GPU info
#!nvidia-smi
```

```
➡ processor      : 0
  vendor_id      : GenuineIntel
  cpu family     : 6
  model         : 85
  model name     : Intel(R) Xeon(R) CPU @ 2.00GHz
  stepping      : 3
  microcode     : 0xffffffff
  cpu MHz       : 2000.192
  cache size    : 39424 KB
  physical id   : 0
```

```

siblings      : 2
core id       : 0
cpu cores     : 1
apicid        : 0
initial apicid : 0
fpu           : yes
fpu_exception : yes
cpuid level   : 13
wp            : yes
flags         : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ss ht syscall n
bugs          : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds swapgs taa mmio_stale_data retbleed bhi
bogomips      : 4000.38
clflush size  : 64
cache_alignment : 64
address sizes  : 46 bits physical, 48 bits virtual
power management:

processor      : 1
vendor_id      : GenuineIntel
cpu family     : 6
model          : 85
model name     : Intel(R) Xeon(R) CPU @ 2.00GHz
stepping       : 3
microcode      : 0xffffffff
cpu MHz        : 2000.192
cache size     : 39424 KB
physical id    : 0
siblings       : 2
core id        : 0
cpu cores      : 1
apicid         : 1
initial apicid : 1
fpu            : yes
fpu_exception  : yes
cpuid level    : 13
wp             : yes
flags          : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ss ht syscall n
bugs           : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds swapgs taa mmio_stale_data retbleed bhi
bogomips       : 4000.38
clflush size   : 64
cache_alignment : 64
address sizes   : 46 bits physical, 48 bits virtual
power management:

```

✓ DOWNLOAD AND CHECK THE DATASET [3 pts]

Data must be located in a ./data directory at the same level as this notebook

#IF YOU USE Google COLAB, you can mount your Google Drive:

```
from google.colab import drive
drive.mount('/content/drive')
```

➞ Mounted at /content/drive

#Create a 'data/' directory, put the numpy files and load. With Colab case, create 'data' here:

```
X_cifar4=np.load("/content/drive/MyDrive/Colab Notebooks/data/CIFAR4_X.npy")
Y_cifar4=np.load("/content/drive/MyDrive/Colab Notebooks/data/CIFAR4_Y.npy")
print(np.shape(X_cifar4))
print(np.shape(Y_cifar4))
nb_labels=4
```

➞ (24000, 32, 32, 3)
(24000,)

[QUESTION] Display some images from X_cifar4 with the corresponding label

NB: Since the pictures are very small, use `plt.figure(figsize=(2, 2), dpi=80)` before `plt.imshow()` to display something "watchable".

```
j = 11995 # Starting index
for i in range(j, j + 9): # Display 9 images sequentially
    plt.figure(figsize=(2, 2), dpi=80) # Set figure size for each image
    plt.imshow(X_cifar4[i]) # Display the image
    plt.title(f"Label: {Y_cifar4[i]}") # Add the label as the title
    plt.axis('off') # Remove axis for clarity
    plt.show() # Show the image
```



Label: 1



Label: 1



Label: 1



Label: 1



Label: 1



Label: 2



Label: 2



Label: 2



Label: 2



A pixel is an UINT8 value, so in $[0;255]$. We will normalize data in $[0,1]$:

```
print("First pixel (r,g,b) of the first image:", X_cifar4[0,0,0,:])
X = X_cifar4/255.0
print("Now in [0,1] ==> ", X[0,0,0,:])
```

```
↵ First pixel (r,g,b) of the first image: [164 206 84]
Now in [0,1] ==> [0.64313725 0.80784314 0.32941176]
```

For the labels **Y**, usually, we prefer to process "one-hot encodings" i.e., a vector with '0' everywhere except for the corresponding label where you have '1'.

Example : let's say you have 4 labels and 10 training data with the original Y:

```
Y=[0,1,1,3,3,3,2,2,1,0]
shape(Y)=[10,]
```

Then the "one-hot encoding" version of Y will be:

```
Y_onehot=
[1,0,0,0;
 0,1,0,0;
 0,1,0,0;
 0,0,0,1;
 0,0,0,1;
 0,0,0,1;
 0,0,1,0;
 0,0,1,0;
 0,1,0,0;
 1,0,0,0]
shape(Y_onehot)=[10,4]
```

[QUESTION] Use the tensorflow method [tf.keras.utils.to_categorical\(\)](#) to transform your Y_cifar4 into Y and check the shape of your new Y.

```
import keras
#Your code here
```

```
Y = keras.utils.to_categorical(Y_cifar4, num_classes=4)
print(Y)
```

```
↵ [[1. 0. 0. 0.]
 [1. 0. 0. 0.]
 [1. 0. 0. 0.]
 ...
 [0. 0. 0. 1.]
 [0. 0. 0. 1.]
 [0. 0. 0. 1.]]
```

✓ CREATE A TRAINING/VALIDATION/TEST dataset [2 pt]

NB: This step is similar to the supervised section in TP1.

[QUESTION] Use the SKLEARN method [train_test_split](#) in order to create:

- a TRAIN set (X_train,Y_train) *[advice: use 70% of the whole dataset]*
- a VALIDATION set (X_val, Y_val) *[advice: 15%]*
- a TEST set (X_test, Y_test) *[advice: 15%]*

Print the shape of the 3 datasets.

The TRAIN and VALIDATION sets will be used at training time. The TEST set will be used, after the training, at the inference time.

NOTE that we could only use a single VALIDATION/TEST set for both the training and the inference stages.

```
#Your code here
#CREATE A TRAINING/ VALIDATION/TEST dataset
from sklearn.model_selection import train_test_split

# Split the dataset into training and test sets
X_train, X_test, Y_train, Y_test= train_test_split(X_cifar4, Y, train_size=0.7, random_state=42)
X_val, X_test, Y_val, Y_test= train_test_split(X_test, Y_test, train_size=0.5, random_state=42)

# Check the sizes of the resulting datasets
print(f"Size of X_train: {X_train.shape}")
print(f"Size of Y_train: {Y_train.shape}")
```

```
print(f"Size of X_test: {X_test.shape}")
print(f"Size of Y_test: {Y_test.shape}")
print(f"Size of X_val: {X_val.shape}")
print(f"Size of Y_val: {Y_val.shape}")
```

```
↗ Size of X_train: (16800, 32, 32, 3)
Size of Y_train: (16800, 4)
Size of X_test: (3600, 32, 32, 3)
Size of Y_test: (3600, 4)
Size of X_val: (3600, 32, 32, 3)
Size of Y_val: (3600, 4)
```

[QUESTION] Why the train/val/test split is important in our case? What is the main interest of the validation data at training time?

[ANSWER] CIFAR-4 is a subset of CIFAR-10 with only 4 classes and limited training examples (24,000 pictures). Small datasets are prone to overfitting: The validation set helps detect overfitting and fine-tune the model. The test set ensures unbiased performance evaluation. And splitting the data ensures the model learns effectively while maintaining a reliable measure of generalization and performance.

As for the validation data, it acts as a proxy for unseen data, enabling:

- Model Monitoring: Track the model's generalization ability during training.
- Early Stopping: Stop training if the validation performance stops improving, preventing overfitting.
- Hyperparameter Tuning: Optimize hyperparameters without contaminating the test set.

Without a validation set, you risk overfitting to the training data and cannot effectively assess generalization during training. This split ensures reliable model evaluation and performance improvement.

✓ MULTI-LAYER PERCEPTRON (MLP) MODEL [12 pts]

✓ Analyse of a MLP code [6pts]

It's time to build our first deep neural network...

Below, we provide a small code **THAT IS NOT WORKING**, because something are missing or wrong.

The model is a very simple MLP with only one hidden layer (20 neurons). So we have :

```
x --> "input" layer ==> hidden layer ==> output layer --> ŷ
```

Have a look at the code, and understand how it works.

```
# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer
x = tf.keras.layers.Flatten()(inputs) #cf. question below...
x = tf.keras.layers.Dense(20, activation='relu')(x) #a first hidden layer with 20 neurons
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_mlp_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_mlp_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_mlp_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
...
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
...

my_mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
...
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
...

nb_epochs=4
batch_size=100
training_history = my_mlp_model.fit(X_train,Y_train,
                                    validation_data=(X_val, Y_val),
                                    epochs=nb_epochs,
```

batch_size=batch_size)

```
#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_mlp_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_mlp_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)
```

Model: "my_mlp_model"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 32, 32, 3)	0
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 20)	61,460
dense_1 (Dense)	(None, 4)	84

Total params: 61,544 (240.41 KB)

Trainable params: 61,544 (240.41 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/4

168/168 ————— 3s 6ms/step - accuracy: 0.2533 - loss: 59.8158 - val_accuracy: 0.2406 - val_loss: 1.3865

Epoch 2/4

168/168 ————— 0s 2ms/step - accuracy: 0.2536 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3865

Epoch 3/4

168/168 ————— 1s 2ms/step - accuracy: 0.2510 - loss: 1.3864 - val_accuracy: 0.2447 - val_loss: 1.3865

Epoch 4/4

168/168 ————— 1s 2ms/step - accuracy: 0.2525 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3864

168/168 ————— 0s 1ms/step - accuracy: 0.2503 - loss: 1.3862

36/36 ————— 0s 2ms/step - accuracy: 0.2509 - loss: 1.3863

Performance on the TRAIN set, ACCURACY= 0.2512499988079071

Performance on the TEST set, ACCURACY= 0.24944444000720978

[QUESTION]

- (1) What is the activation function for the hidden layer?
- (2) What is the most used activation function in deep learning?
- (3) How many time an image sample will be used during the training?
- (4) How many training iterations (i.e., params update) in total will be processed?

[ANSWER]

(1) The activation function for the hidden layer is `relu()`.

(2) The the most used activation function in deep learning is `relu()` because ,even with its cons , it is a great way to avoid vanishing gradient and it allows a sparse activation of our network.

(3) Because this deep learning algorithm runs on 4 epochs, each image sample will be used 4 times during the training.

(4) The training dataset have 70% of image samples of our ciphar4 dataset, which means 16800 pictures. Furthermore, the MLP algorithm above is doing a mini-batch Stochastic Gradient Descent with batches of size 100, which means the training dataset have 168 batches. And this algorithm runs on 4 epochs, so in total $672 (=4 \times 24000 \times 0.7 / 100)$ training iterations will be processed.

[QUESTION]

- (1) Regarding the previous questions, copy and change the code above to make it runs on **10 EPOCHS**. Comment the resulting performance of the model.
- (2) For our MLP, why do we need this line `x = tf.keras.layers.Flatten()(inputs) ?`
- (3) Give the calculation of the number of trainable parameters

#Your code here

(1) DEFINE THE ARCHITECTURE OF MY MODEL

#first, I define all the layers and the way they are connected

inputs = tf.keras.Input(shape=(32,32,3)) #my input layer

x = tf.keras.layers.Flatten()(inputs) #cf. question below...

x = tf.keras.layers.Dense(20, activation='relu')(x) #a first hidden layer with 20 neurons

outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer

#Then, I define my model with the input layer, the output layer and a name

my_mlp_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_mlp_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS

my_mlp_model.summary()

(2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:

...

(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']

```

(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=10
batch_size=100
training_history = my_mlp_model.fit(X_train,Y_train,
                                    validation_data=(X_val, Y_val),
                                    epochs=nb_epochs,
                                    batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_mlp_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_mlp_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

```

→ Model: "my_mlp_model"

Layer (type)	Output Shape	Param #
input_layer_6 (InputLayer)	(None, 32, 32, 3)	0
flatten_4 (Flatten)	(None, 3072)	0
dense_12 (Dense)	(None, 20)	61,460
dense_13 (Dense)	(None, 4)	84

Total params: 61,544 (240.41 KB)

Trainable params: 61,544 (240.41 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

168/168 ————— 2s 5ms/step - accuracy: 0.2566 - loss: 76.4101 - val_accuracy: 0.2406 - val_loss: 1.3866

Epoch 2/10

168/168 ————— 0s 2ms/step - accuracy: 0.2492 - loss: 1.3864 - val_accuracy: 0.2406 - val_loss: 1.3865

Epoch 3/10

168/168 ————— 1s 2ms/step - accuracy: 0.2364 - loss: 1.3864 - val_accuracy: 0.2406 - val_loss: 1.3865

Epoch 4/10

168/168 ————— 1s 2ms/step - accuracy: 0.2522 - loss: 1.3863 - val_accuracy: 0.2406 - val_loss: 1.3865

Epoch 5/10

168/168 ————— 1s 2ms/step - accuracy: 0.2487 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3864

Epoch 6/10

168/168 ————— 1s 2ms/step - accuracy: 0.2505 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3864

Epoch 7/10

168/168 ————— 0s 2ms/step - accuracy: 0.2505 - loss: 1.3863 - val_accuracy: 0.2406 - val_loss: 1.3864

Epoch 8/10

168/168 ————— 1s 2ms/step - accuracy: 0.2486 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3864

Epoch 9/10

168/168 ————— 0s 2ms/step - accuracy: 0.2539 - loss: 1.3863 - val_accuracy: 0.2447 - val_loss: 1.3865

Epoch 10/10

168/168 ————— 1s 2ms/step - accuracy: 0.2563 - loss: 1.3862 - val_accuracy: 0.2406 - val_loss: 1.3864

168/168 ————— 0s 1ms/step - accuracy: 0.2589 - loss: 1.3862

36/36 ————— 0s 2ms/step - accuracy: 0.2522 - loss: 1.3863

Performance on the TRAIN set, ACCURACY= 0.25065475702285767

Performance on the TEST set, ACCURACY= 0.25638890766418457

[ANSWER] Your answer here

(1) Through all of our 10 epochs, we have an accuracy around 25%. However this classification task has only 4 labels equally distributed in the dataset which means that our model is as good as just choosing randomly a picture. And this accuracy of 25% is the case for both the train and the test set which means that our model is clearly **underfitting**.

(2) The line `x = tf.keras.layers.Flatten()(inputs)` is essential for preparing the input data for the Dense (fully connected) layers in the MLP architecture.

The input data has the shape (32, 32, 3), representing an RGB image of size 32×32 with 3 color channels. However dense Layers Require 1D Input. This is why we have a Flatten layer (`x = tf.keras.layers.Flatten()(inputs)`) that converts the 3D input tensor (height, width, channels) into a 1D vector.

(3) We have 2 layers, the first one with 20 neurons and the second one with 4 neurons. Which means that we have 84 $(=(20+1) \times 4)$ trainable parameters if we count the biases.


It is always good to have a look on the training curves. That is the role of the **"training_history"** object that we defined in the code as the output of the **fit** method. In this object, we collect all the loss and metric values after each epoch.

You can use the following **method** to display the train/val curves (loss and accuracy):

```
def display_training_curves(training_history):
    # display loss
    plt.plot(training_history.history['loss'], label='loss')
    plt.plot(training_history.history['val_loss'], label='val_loss')
    plt.xlabel("epoch")
    plt.legend(bbox_to_anchor=(1.05, 1), loc=2)
    plt.show()
    #display accuracy
    plt.plot(training_history.history['accuracy'], label='acc')
    plt.plot(training_history.history['val_accuracy'], label='val_acc')
    plt.xlabel("epoch")
    plt.legend(bbox_to_anchor=(1.05, 1), loc=2)
    plt.show()
```

[QUESTION] Display the training curves from your MLP model.

```
display_training_curves(training_history)
```



Improvements and overfitting [6 pts]

[QUESTION] Use the previous MLP code to build your own model and try to reach a better accuracy performance (e.g., above/around 60%, both train and test)


```
#Your code here
```

```
# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer
x = tf.keras.layers.Flatten()(inputs) #cf. question below...
x = tf.keras.layers.Dense(1024, activation='relu')(x) #a first hidden layer with 1024 neurons
x = tf.keras.layers.Dense(768, activation='relu')(x) #a second hidden layer with 768 neurons
x = tf.keras.layers.Dense(512, activation='relu')(x) #a second hidden layer with 512 neurons
x = tf.keras.layers.Dense(256, activation='relu')(x) #a second hidden layer with 256 neurons
x = tf.keras.layers.Dense(128, activation='relu')(x) #a second hidden layer with 128 neurons
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_mlp_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_mlp_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_mlp_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''
nb_epochs=22
batch_size=100
training_history = my_mlp_model.fit(X_train,Y_train,
                                   validation_data=(X_val, Y_val),
                                   epochs=nb_epochs,
                                   batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_mlp_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_mlp_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)
```

Model: "my_mlp_model"

Layer (type)	Output Shape	Param #
input_layer_7 (InputLayer)	(None, 32, 32, 3)	0
flatten_5 (Flatten)	(None, 3072)	0
dense_14 (Dense)	(None, 1024)	3,146,752
dense_15 (Dense)	(None, 768)	787,200
dense_16 (Dense)	(None, 512)	393,728
dense_17 (Dense)	(None, 256)	131,328
dense_18 (Dense)	(None, 128)	32,896
dense_19 (Dense)	(None, 4)	516

Total params: 4,492,420 (17.14 MB)

Trainable params: 4,492,420 (17.14 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/22

168/168 ————— 4s 8ms/step - accuracy: 0.2951 - loss: 98.4515 - val_accuracy: 0.4431 - val_loss: 1.2845

Epoch 2/22

168/168 ————— 1s 3ms/step - accuracy: 0.4809 - loss: 1.2057 - val_accuracy: 0.5278 - val_loss: 1.1410

Epoch 3/22

168/168 ————— 1s 3ms/step - accuracy: 0.5181 - loss: 1.1331 - val_accuracy: 0.5417 - val_loss: 1.0966

Epoch 4/22

168/168 ————— 1s 3ms/step - accuracy: 0.5469 - loss: 1.0766 - val_accuracy: 0.5264 - val_loss: 1.1308

Epoch 5/22

168/168 ————— 1s 3ms/step - accuracy: 0.5730 - loss: 1.0358 - val_accuracy: 0.5669 - val_loss: 1.0396

Epoch 6/22

168/168 ————— 1s 3ms/step - accuracy: 0.5634 - loss: 1.0461 - val_accuracy: 0.5669 - val_loss: 1.0457

Epoch 7/22

168/168 ————— 1s 3ms/step - accuracy: 0.5721 - loss: 1.0150 - val_accuracy: 0.5322 - val_loss: 1.1097

Epoch 8/22

168/168 ————— 1s 4ms/step - accuracy: 0.5710 - loss: 1.0289 - val_accuracy: 0.5853 - val_loss: 1.0232

Epoch 9/22

168/168 ————— 1s 4ms/step - accuracy: 0.6026 - loss: 0.9538 - val_accuracy: 0.5783 - val_loss: 1.0148

Epoch 10/22

168/168 ————— 1s 5ms/step - accuracy: 0.6018 - loss: 0.9657 - val_accuracy: 0.5736 - val_loss: 1.0351

Epoch 11/22

168/168 ————— 1s 4ms/step - accuracy: 0.5959 - loss: 0.9706 - val_accuracy: 0.5908 - val_loss: 0.9961

Epoch 12/22

168/168 ————— 1s 3ms/step - accuracy: 0.6139 - loss: 0.9473 - val_accuracy: 0.5869 - val_loss: 1.0101

Epoch 13/22

168/168 ————— 1s 3ms/step - accuracy: 0.6192 - loss: 0.9190 - val_accuracy: 0.5556 - val_loss: 1.0578

Epoch 14/22

168/168 ————— 1s 3ms/step - accuracy: 0.5969 - loss: 0.9774 - val_accuracy: 0.5881 - val_loss: 1.0106

Epoch 15/22

168/168 ————— 1s 3ms/step - accuracy: 0.6263 - loss: 0.9128 - val_accuracy: 0.5628 - val_loss: 1.0635

Epoch 16/22

168/168 ————— 1s 4ms/step - accuracy: 0.6236 - loss: 0.9174 - val_accuracy: 0.6083 - val_loss: 0.9879

Epoch 17/22

168/168 ————— 1s 3ms/step - accuracy: 0.6240 - loss: 0.9064 - val_accuracy: 0.6003 - val_loss: 1.0023

Epoch 18/22

168/168 ————— 1s 3ms/step - accuracy: 0.6222 - loss: 0.9166 - val_accuracy: 0.5853 - val_loss: 1.0323

Epoch 19/22

168/168 ————— 1s 3ms/step - accuracy: 0.6211 - loss: 0.9128 - val_accuracy: 0.5472 - val_loss: 1.0855

Epoch 20/22

168/168 ————— 1s 4ms/step - accuracy: 0.6299 - loss: 0.9156 - val_accuracy: 0.5975 - val_loss: 1.0084

Epoch 21/22

168/168 ————— 1s 3ms/step - accuracy: 0.6308 - loss: 0.9071 - val_accuracy: 0.6200 - val_loss: 0.9768

Epoch 22/22

168/168 ————— 1s 3ms/step - accuracy: 0.6438 - loss: 0.8654 - val_accuracy: 0.5719 - val_loss: 1.0340

168/168 ————— 0s 2ms/step - accuracy: 0.6338 - loss: 0.8785

36/36 ————— 0s 2ms/step - accuracy: 0.5597 - loss: 1.0570

Performance on the TRAIN set, ACCURACY= 0.6311904788017273

Performance on the TEST set, ACCURACY= 0.5616666674613953

[QUESTION] For illustration/educational purpose, use the code of your last model and adapt it so that your model clearly **OVERFITS**.

We need to see the overfitting issue on the training curve!

(Think about the reasons of overfitting?)

#Your code here

(1) DEFINE THE ARCHITECTURE OF MY MODEL

#first, I define all the layers and the way they are connected

inputs = tf.keras.Input(shape=(32,32,3)) #my input layer

x = tf.keras.layers.Flatten()(inputs) #cf. question below...

x = tf.keras.layers.Dense(1024, activation='relu')(x) #a first hidden layer with 1024 neurons

x = tf.keras.layers.Dense(768, activation='relu')(x) #a second hidden layer with 768 neurons

x = tf.keras.layers.Dense(512, activation='relu')(x) #a second hidden layer with 512 neurons

```

x = tf.keras.layers.Dense(256, activation='relu')(x) #a second hidden layer with 256 neurons
x = tf.keras.layers.Dense(128, activation='relu')(x) #a second hidden layer with 128 neurons
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_mlp_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_mlp_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_mlp_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=60
batch_size=100
training_history = my_mlp_model.fit(X_train,Y_train,
                                   validation_data=(X_val, Y_val),
                                   epochs=nb_epochs,
                                   batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_mlp_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_mlp_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

display_training_curves(training_history)

```

Model: "my_mlp_model"

Layer (type)	Output Shape	Param #
input_layer_8 (InputLayer)	(None, 32, 32, 3)	0
flatten_6 (Flatten)	(None, 3072)	0
dense_20 (Dense)	(None, 1024)	3,146,752
dense_21 (Dense)	(None, 768)	787,200
dense_22 (Dense)	(None, 512)	393,728
dense_23 (Dense)	(None, 256)	131,328
dense_24 (Dense)	(None, 128)	32,896
dense_25 (Dense)	(None, 4)	516

Total params: 4,492,420 (17.14 MB)

Trainable params: 4,492,420 (17.14 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/60

168/168 ————— 4s 8ms/step - accuracy: 0.2847 - loss: 119.3022 - val_accuracy: 0.3883 - val_loss: 1.4112

Epoch 2/60

168/168 ————— 1s 5ms/step - accuracy: 0.4485 - loss: 1.2745 - val_accuracy: 0.5039 - val_loss: 1.1479

Epoch 3/60

168/168 ————— 1s 5ms/step - accuracy: 0.5182 - loss: 1.1309 - val_accuracy: 0.4956 - val_loss: 1.1697

Epoch 4/60

168/168 ————— 1s 5ms/step - accuracy: 0.5228 - loss: 1.1212 - val_accuracy: 0.4847 - val_loss: 1.2066

Epoch 5/60

168/168 ————— 1s 4ms/step - accuracy: 0.5377 - loss: 1.0870 - val_accuracy: 0.5297 - val_loss: 1.1188

Epoch 6/60

168/168 ————— 1s 3ms/step - accuracy: 0.5569 - loss: 1.0568 - val_accuracy: 0.5647 - val_loss: 1.0615

Epoch 7/60

168/168 ————— 1s 3ms/step - accuracy: 0.5742 - loss: 1.0196 - val_accuracy: 0.5728 - val_loss: 1.0252

Epoch 8/60

168/168 ————— 1s 3ms/step - accuracy: 0.5850 - loss: 0.9861 - val_accuracy: 0.5294 - val_loss: 1.1151

Epoch 9/60

168/168 ————— 1s 3ms/step - accuracy: 0.5866 - loss: 0.9986 - val_accuracy: 0.5706 - val_loss: 1.0250

Epoch 10/60

168/168 ————— 1s 3ms/step - accuracy: 0.5935 - loss: 0.9896 - val_accuracy: 0.5814 - val_loss: 1.0342

Epoch 11/60

168/168 ————— 1s 3ms/step - accuracy: 0.5902 - loss: 0.9808 - val_accuracy: 0.5986 - val_loss: 1.0103

Epoch 12/60

168/168 ————— 1s 3ms/step - accuracy: 0.6075 - loss: 0.9495 - val_accuracy: 0.5556 - val_loss: 1.0882

Epoch 13/60

168/168 ————— 1s 3ms/step - accuracy: 0.6053 - loss: 0.9440 - val_accuracy: 0.5439 - val_loss: 1.1298

Epoch 14/60

168/168 ————— 1s 3ms/step - accuracy: 0.6108 - loss: 0.9511 - val_accuracy: 0.5831 - val_loss: 1.0380

Epoch 15/60

168/168 ————— 1s 3ms/step - accuracy: 0.6116 - loss: 0.9466 - val_accuracy: 0.5428 - val_loss: 1.0931

Epoch 16/60

168/168 ————— 1s 3ms/step - accuracy: 0.6128 - loss: 0.9419 - val_accuracy: 0.5697 - val_loss: 1.0643

Epoch 17/60

168/168 ————— 1s 3ms/step - accuracy: 0.6226 - loss: 0.9295 - val_accuracy: 0.5550 - val_loss: 1.1019

Epoch 18/60

168/168 ————— 1s 3ms/step - accuracy: 0.6077 - loss: 0.9410 - val_accuracy: 0.5811 - val_loss: 1.0243

Epoch 19/60

168/168 ————— 1s 3ms/step - accuracy: 0.6377 - loss: 0.8977 - val_accuracy: 0.5661 - val_loss: 1.1208

Epoch 20/60

168/168 ————— 1s 4ms/step - accuracy: 0.6346 - loss: 0.8932 - val_accuracy: 0.5544 - val_loss: 1.1217

Epoch 21/60

168/168 ————— 1s 5ms/step - accuracy: 0.6388 - loss: 0.8934 - val_accuracy: 0.5875 - val_loss: 1.0414

Epoch 22/60

168/168 ————— 1s 5ms/step - accuracy: 0.6408 - loss: 0.8846 - val_accuracy: 0.6014 - val_loss: 0.9890

Epoch 23/60

168/168 ————— 1s 5ms/step - accuracy: 0.6357 - loss: 0.8809 - val_accuracy: 0.5950 - val_loss: 1.0110

Epoch 24/60

168/168 ————— 1s 4ms/step - accuracy: 0.6357 - loss: 0.8792 - val_accuracy: 0.6017 - val_loss: 0.9995

Epoch 25/60

168/168 ————— 1s 3ms/step - accuracy: 0.6500 - loss: 0.8552 - val_accuracy: 0.5986 - val_loss: 0.9936

Epoch 26/60

168/168 ————— 1s 3ms/step - accuracy: 0.6582 - loss: 0.8403 - val_accuracy: 0.5514 - val_loss: 1.1204

Epoch 27/60

168/168 ————— 1s 4ms/step - accuracy: 0.6613 - loss: 0.8385 - val_accuracy: 0.5456 - val_loss: 1.1353

Epoch 28/60

168/168 ————— 1s 3ms/step - accuracy: 0.6317 - loss: 0.8884 - val_accuracy: 0.6131 - val_loss: 1.0070

Epoch 29/60

168/168 ————— 1s 3ms/step - accuracy: 0.6627 - loss: 0.8338 - val_accuracy: 0.5783 - val_loss: 1.0610

Epoch 30/60

168/168 ————— 1s 3ms/step - accuracy: 0.6595 - loss: 0.8340 - val_accuracy: 0.5761 - val_loss: 1.1023

Epoch 31/60

168/168 ————— 1s 3ms/step - accuracy: 0.6667 - loss: 0.8288 - val_accuracy: 0.6000 - val_loss: 1.0120

Epoch 32/60

168/168 ————— 1s 3ms/step - accuracy: 0.6672 - loss: 0.8116 - val_accuracy: 0.6106 - val_loss: 0.9890



Epoch 33/60


168/168 ————— 1s 3ms/step - accuracy: 0.6731 - loss: 0.8134 - val_accuracy: 0.5892 - val_loss: 1.0427

Epoch 34/60

168/168 ————— 1s 3ms/step - accuracy: 0.6680 - loss: 0.8231 - val_accuracy: 0.6072 - val_loss: 1.0132

```
Epoch 35/60
168/168 1s 3ms/step - accuracy: 0.6673 - loss: 0.8223 - val_accuracy: 0.5986 - val_loss: 1.0118
Epoch 36/60
168/168 1s 3ms/step - accuracy: 0.6678 - loss: 0.8194 - val_accuracy: 0.5839 - val_loss: 1.0335
Epoch 37/60
168/168 1s 3ms/step - accuracy: 0.6913 - loss: 0.7650 - val_accuracy: 0.6072 - val_loss: 0.9784
Epoch 38/60
168/168 1s 3ms/step - accuracy: 0.6995 - loss: 0.7473 - val_accuracy: 0.6094 - val_loss: 1.0100
Epoch 39/60
168/168 1s 4ms/step - accuracy: 0.6953 - loss: 0.7567 - val_accuracy: 0.6150 - val_loss: 1.0258
Epoch 40/60
168/168 1s 5ms/step - accuracy: 0.7031 - loss: 0.7548 - val_accuracy: 0.5919 - val_loss: 1.0451
Epoch 41/60
168/168 1s 5ms/step - accuracy: 0.6857 - loss: 0.7826 - val_accuracy: 0.6014 - val_loss: 1.0454
Epoch 42/60
168/168 1s 5ms/step - accuracy: 0.7050 - loss: 0.7300 - val_accuracy: 0.5875 - val_loss: 1.0545
Epoch 43/60
168/168 1s 4ms/step - accuracy: 0.7100 - loss: 0.7265 - val_accuracy: 0.5772 - val_loss: 1.1260
Epoch 44/60
168/168 1s 3ms/step - accuracy: 0.7026 - loss: 0.7420 - val_accuracy: 0.6072 - val_loss: 1.0386
Epoch 45/60
168/168 1s 3ms/step - accuracy: 0.6989 - loss: 0.7441 - val_accuracy: 0.6183 - val_loss: 1.0064
Epoch 46/60
168/168 1s 3ms/step - accuracy: 0.7189 - loss: 0.7108 - val_accuracy: 0.6089 - val_loss: 1.0161
Epoch 47/60
168/168 1s 3ms/step - accuracy: 0.7244 - loss: 0.6989 - val_accuracy: 0.6108 - val_loss: 1.0338
Epoch 48/60
168/168 1s 3ms/step - accuracy: 0.7158 - loss: 0.7097 - val_accuracy: 0.5972 - val_loss: 1.0899
Epoch 49/60
168/168 1s 3ms/step - accuracy: 0.7179 - loss: 0.7129 - val_accuracy: 0.6064 - val_loss: 1.0159
Epoch 50/60
168/168 1s 3ms/step - accuracy: 0.7252 - loss: 0.6833 - val_accuracy: 0.6194 - val_loss: 1.0414
Epoch 51/60
168/168 1s 3ms/step - accuracy: 0.7149 - loss: 0.7082 - val_accuracy: 0.6192 - val_loss: 1.0478
Epoch 52/60
168/168 1s 3ms/step - accuracy: 0.7293 - loss: 0.6809 - val_accuracy: 0.6178 - val_loss: 1.0380
Epoch 53/60
168/168 1s 4ms/step - accuracy: 0.7309 - loss: 0.6757 - val_accuracy: 0.6136 - val_loss: 1.0461
Epoch 54/60
168/168 1s 3ms/step - accuracy: 0.7376 - loss: 0.6584 - val_accuracy: 0.6194 - val_loss: 1.0126
Epoch 55/60
168/168 1s 3ms/step - accuracy: 0.7336 - loss: 0.6770 - val_accuracy: 0.6094 - val_loss: 1.0848
Epoch 56/60
168/168 1s 3ms/step - accuracy: 0.7454 - loss: 0.6537 - val_accuracy: 0.6292 - val_loss: 1.0800
Epoch 57/60
168/168 1s 3ms/step - accuracy: 0.7428 - loss: 0.6455 - val_accuracy: 0.6019 - val_loss: 1.0769
Epoch 58/60
168/168 1s 3ms/step - accuracy: 0.7446 - loss: 0.6442 - val_accuracy: 0.5956 - val_loss: 1.1276
Epoch 59/60
168/168 1s 4ms/step - accuracy: 0.7339 - loss: 0.6679 - val_accuracy: 0.6164 - val_loss: 1.2128
Epoch 60/60
168/168 1s 5ms/step - accuracy: 0.7638 - loss: 0.6027 - val_accuracy: 0.6264 - val_loss: 1.0702
168/168 0s 3ms/step - accuracy: 0.7637 - loss: 0.6055
36/36 0s 2ms/step - accuracy: 0.6150 - loss: 1.0940
Performance on the TRAIN set, ACCURACY= 0.7672619223594666
Performance on the TEST set, ACCURACY= 0.6113888621330261
```





[QUESTION] Try to find an optimal architecture without overfitting by using a regularization (or other) technique of your choice. The goal is to have a performance > 60% without overfitting after **20** epochs.

#Your code here

```
# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer
x = tf.keras.layers.Flatten()(inputs) #cf. question below...

#x = tf.keras.layers.Dense(1024, activation='relu',kernel_regularizer=tf.keras.regularizers.l2(0.01))(x) does not work when doing it on
#x = tf.keras.layers.BatchNormalization()(x) : ultra overfitting (training_acc : 95% and test_accuracy : 60%)
#x = tf.keras.layers.Dropout(0.1)(x) Dropping 10% of neurons on the first layer only, allows to eliminates an important part of our pro

x = tf.keras.layers.Dense(1024, activation='relu')(x) #a first hidden layer with 1024 neurons
x = tf.keras.layers.Dropout(0.1)(x) #Drop 10% of neurons on the first layer
x = tf.keras.layers.Dense(768, activation='relu')(x) #a second hidden layer with 768 neurons
x = tf.keras.layers.Dense(512, activation='relu')(x) #a second hidden layer with 512 neurons
x = tf.keras.layers.Dense(256, activation='relu')(x) #a second hidden layer with 256 neurons
x = tf.keras.layers.Dense(128, activation='relu')(x) #a second hidden layer with 128 neurons
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_mlp_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_mlp_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_mlp_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=60
batch_size=100
training_history = my_mlp_model.fit(X_train,Y_train,
                                   validation_data=(X_val, Y_val),
                                   epochs=nb_epochs,
                                   batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_mlp_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_mlp_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)
display_training_curves(training_history)
```

Model: "my_mlp_model"

Layer (type)	Output Shape	Param #
input_layer_11 (InputLayer)	(None, 32, 32, 3)	0
flatten_9 (Flatten)	(None, 3072)	0
dense_38 (Dense)	(None, 1024)	3,146,752
dropout_2 (Dropout)	(None, 1024)	0
dense_39 (Dense)	(None, 768)	787,200
dense_40 (Dense)	(None, 512)	393,728
dense_41 (Dense)	(None, 256)	131,328
dense_42 (Dense)	(None, 128)	32,896
dense_43 (Dense)	(None, 4)	516

Total params: 4,492,420 (17.14 MB)

Trainable params: 4,492,420 (17.14 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/60

168/168 ————— 5s 11ms/step - accuracy: 0.2687 - loss: 119.5292 - val_accuracy: 0.3694 - val_loss: 1.3006

Epoch 2/60

168/168 ————— 2s 4ms/step - accuracy: 0.3617 - loss: 1.4203 - val_accuracy: 0.4206 - val_loss: 1.2740

Epoch 3/60

168/168 ————— 1s 3ms/step - accuracy: 0.4427 - loss: 1.2367 - val_accuracy: 0.5222 - val_loss: 1.1156

Epoch 4/60

168/168 ————— 1s 3ms/step - accuracy: 0.5040 - loss: 1.1457 - val_accuracy: 0.5428 - val_loss: 1.0827

Epoch 5/60

168/168 ————— 1s 3ms/step - accuracy: 0.5206 - loss: 1.1149 - val_accuracy: 0.5633 - val_loss: 1.0586

Epoch 6/60

168/168 ————— 1s 3ms/step - accuracy: 0.5273 - loss: 1.1056 - val_accuracy: 0.5364 - val_loss: 1.0980

Epoch 7/60

168/168 ————— 1s 3ms/step - accuracy: 0.5273 - loss: 1.1056 - val_accuracy: 0.5597 - val_loss: 1.0470

Epoch 8/60

168/168 ————— 1s 3ms/step - accuracy: 0.5454 - loss: 1.0775 - val_accuracy: 0.5572 - val_loss: 1.0689

Epoch 9/60

168/168 ————— 1s 3ms/step - accuracy: 0.5352 - loss: 1.0753 - val_accuracy: 0.5544 - val_loss: 1.0647

Epoch 10/60

168/168 ————— 1s 3ms/step - accuracy: 0.5474 - loss: 1.0588 - val_accuracy: 0.5475 - val_loss: 1.0769

Epoch 11/60

168/168 ————— 1s 3ms/step - accuracy: 0.5414 - loss: 1.0716 - val_accuracy: 0.5425 - val_loss: 1.0772

Epoch 12/60

168/168 ————— 1s 3ms/step - accuracy: 0.5524 - loss: 1.0403 - val_accuracy: 0.5708 - val_loss: 1.0368

Epoch 13/60

168/168 ————— 1s 3ms/step - accuracy: 0.5559 - loss: 1.0546 - val_accuracy: 0.5253 - val_loss: 1.1143

Epoch 14/60

168/168 ————— 1s 3ms/step - accuracy: 0.5446 - loss: 1.0552 - val_accuracy: 0.5681 - val_loss: 1.0334

Epoch 15/60

168/168 ————— 1s 4ms/step - accuracy: 0.5806 - loss: 1.0059 - val_accuracy: 0.5475 - val_loss: 1.0619

Epoch 16/60

168/168 ————— 1s 4ms/step - accuracy: 0.5658 - loss: 1.0288 - val_accuracy: 0.5806 - val_loss: 1.0083

Epoch 17/60

168/168 ————— 1s 5ms/step - accuracy: 0.5819 - loss: 1.0005 - val_accuracy: 0.5647 - val_loss: 1.0294

Epoch 18/60

168/168 ————— 1s 5ms/step - accuracy: 0.5742 - loss: 1.0033 - val_accuracy: 0.5753 - val_loss: 1.0250

Epoch 19/60

168/168 ————— 1s 4ms/step - accuracy: 0.5853 - loss: 0.9875 - val_accuracy: 0.5497 - val_loss: 1.0627

Epoch 20/60

168/168 ————— 1s 3ms/step - accuracy: 0.5844 - loss: 0.9882 - val_accuracy: 0.5575 - val_loss: 1.0409

Epoch 21/60

168/168 ————— 1s 3ms/step - accuracy: 0.5598 - loss: 1.0279 - val_accuracy: 0.5906 - val_loss: 0.9992

Epoch 22/60

168/168 ————— 1s 3ms/step - accuracy: 0.5930 - loss: 0.9764 - val_accuracy: 0.5553 - val_loss: 1.0715

Epoch 23/60

168/168 ————— 1s 3ms/step - accuracy: 0.5830 - loss: 0.9980 - val_accuracy: 0.5881 - val_loss: 1.0117

Epoch 24/60

168/168 ————— 1s 3ms/step - accuracy: 0.5928 - loss: 0.9759 - val_accuracy: 0.5617 - val_loss: 1.0450

Epoch 25/60

168/168 ————— 1s 3ms/step - accuracy: 0.5814 - loss: 0.9914 - val_accuracy: 0.5803 - val_loss: 0.9958

Epoch 26/60

168/168 ————— 1s 3ms/step - accuracy: 0.5968 - loss: 0.9685 - val_accuracy: 0.5750 - val_loss: 1.0047

Epoch 27/60

168/168 ————— 1s 3ms/step - accuracy: 0.5936 - loss: 0.9658 - val_accuracy: 0.5867 - val_loss: 1.0057

Epoch 28/60

168/168 ————— 1s 3ms/step - accuracy: 0.5936 - loss: 0.9670 - val_accuracy: 0.5725 - val_loss: 1.0314

Epoch 29/60

168/168 ————— 1s 3ms/step - accuracy: 0.6178 - loss: 0.9396 - val_accuracy: 0.5936 - val_loss: 0.9883

Epoch 30/60

168/168 ————— 1s 3ms/step - accuracy: 0.6094 - loss: 0.9454 - val_accuracy: 0.5889 - val_loss: 1.0120

Epoch 31/60

168/168 ————— 1s 3ms/step - accuracy: 0.5956 - loss: 0.9559 - val_accuracy: 0.5686 - val_loss: 1.0411



Epoch 32/60


168/168 ————— 1s 3ms/step - accuracy: 0.5960 - loss: 0.9706 - val_accuracy: 0.5908 - val_loss: 1.0185

Epoch 33/60

168/168 ————— 1s 3ms/step - accuracy: 0.6113 - loss: 0.9395 - val_accuracy: 0.5875 - val_loss: 1.0005


```
Epoch 34/60
168/168 1s 3ms/step - accuracy: 0.6115 - loss: 0.9388 - val_accuracy: 0.5997 - val_loss: 0.9695
Epoch 35/60
168/168 1s 3ms/step - accuracy: 0.6116 - loss: 0.9320 - val_accuracy: 0.6033 - val_loss: 0.9685
Epoch 36/60
168/168 1s 5ms/step - accuracy: 0.6237 - loss: 0.9137 - val_accuracy: 0.5800 - val_loss: 1.0091
Epoch 37/60
168/168 1s 5ms/step - accuracy: 0.6089 - loss: 0.9449 - val_accuracy: 0.5808 - val_loss: 1.0166
Epoch 38/60
168/168 1s 5ms/step - accuracy: 0.6084 - loss: 0.9446 - val_accuracy: 0.5942 - val_loss: 0.9789
Epoch 39/60
168/168 1s 3ms/step - accuracy: 0.6248 - loss: 0.9112 - val_accuracy: 0.6006 - val_loss: 0.9828
Epoch 40/60
168/168 1s 3ms/step - accuracy: 0.6110 - loss: 0.9341 - val_accuracy: 0.5983 - val_loss: 0.9754
Epoch 41/60
168/168 1s 3ms/step - accuracy: 0.6233 - loss: 0.9077 - val_accuracy: 0.6031 - val_loss: 1.0095
Epoch 42/60
168/168 1s 3ms/step - accuracy: 0.6334 - loss: 0.8893 - val_accuracy: 0.6100 - val_loss: 0.9674
Epoch 43/60
168/168 1s 3ms/step - accuracy: 0.6338 - loss: 0.8798 - val_accuracy: 0.5936 - val_loss: 0.9816
Epoch 44/60
168/168 1s 3ms/step - accuracy: 0.6288 - loss: 0.8955 - val_accuracy: 0.5814 - val_loss: 1.0120
Epoch 45/60
168/168 1s 3ms/step - accuracy: 0.6368 - loss: 0.8893 - val_accuracy: 0.5997 - val_loss: 0.9915
Epoch 46/60
168/168 1s 3ms/step - accuracy: 0.6355 - loss: 0.8918 - val_accuracy: 0.5800 - val_loss: 1.0147
Epoch 47/60
168/168 1s 3ms/step - accuracy: 0.6281 - loss: 0.8967 - val_accuracy: 0.6044 - val_loss: 0.9728
Epoch 48/60
168/168 1s 3ms/step - accuracy: 0.6340 - loss: 0.8809 - val_accuracy: 0.5897 - val_loss: 1.0087
Epoch 49/60
168/168 1s 3ms/step - accuracy: 0.6145 - loss: 0.9282 - val_accuracy: 0.5981 - val_loss: 1.0048
Epoch 50/60
168/168 1s 4ms/step - accuracy: 0.6370 - loss: 0.8674 - val_accuracy: 0.5939 - val_loss: 1.0208
Epoch 51/60
168/168 1s 3ms/step - accuracy: 0.6396 - loss: 0.8808 - val_accuracy: 0.5989 - val_loss: 0.9951
Epoch 52/60
168/168 1s 3ms/step - accuracy: 0.6436 - loss: 0.8707 - val_accuracy: 0.6114 - val_loss: 0.9733
Epoch 53/60
168/168 1s 3ms/step - accuracy: 0.6424 - loss: 0.8657 - val_accuracy: 0.5894 - val_loss: 0.9974
Epoch 54/60
168/168 1s 4ms/step - accuracy: 0.6419 - loss: 0.8788 - val_accuracy: 0.5997 - val_loss: 0.9708
Epoch 55/60
168/168 1s 4ms/step - accuracy: 0.6360 - loss: 0.8898 - val_accuracy: 0.6100 - val_loss: 0.9790
Epoch 56/60
168/168 1s 5ms/step - accuracy: 0.6500 - loss: 0.8553 - val_accuracy: 0.5903 - val_loss: 1.0230
Epoch 57/60
168/168 1s 4ms/step - accuracy: 0.6467 - loss: 0.8647 - val_accuracy: 0.6081 - val_loss: 0.9667
Epoch 58/60
168/168 1s 3ms/step - accuracy: 0.6463 - loss: 0.8741 - val_accuracy: 0.5981 - val_loss: 0.9855
Epoch 59/60
168/168 1s 3ms/step - accuracy: 0.6439 - loss: 0.8718 - val_accuracy: 0.6061 - val_loss: 0.9689
Epoch 60/60
168/168 1s 3ms/step - accuracy: 0.6432 - loss: 0.8592 - val_accuracy: 0.5981 - val_loss: 0.9956
0s 2ms/step - accuracy: 0.6670 - loss: 0.8248
36/36 0s 2ms/step - accuracy: 0.5836 - loss: 1.0136
Performance on the TRAIN set, ACCURACY= 0.6651190519332886
Performance on the TEST set, ACCURACY= 0.5869444608688354
```





✓ CONVOLUTIONAL NEURAL NETWORK (CNN) [12 pts]

✓ Build a first architecture [6 pts]

MLPs are great but CNNs should work better for our image classification problem...

For that, we will use new layers from TF.KERAS:

- [tf.keras.layers.Conv2D\(\)](#) an example is:

```
l = tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu')(l_input)
```

Here, we ask for 32 convolutional kernels of size [3,3]. By default the stride is set to '1' and the padding is 'valid'.

- [tf.keras.layers.MaxPooling2D\(\)](#) an example is:

```
l = tf.keras.layers.MaxPooling2D()(l_input)
```

By default the stride is set to '2' and the padding is 'valid'.

[QUESTION] With the default parameters of *Conv2D()*, do you expect to have the same shape for the output tensor?

[ANSWER] Not at all, because of the padding='valid' (so no padding added around the input) by default, the output tensor will at least shrink by 1 in height and width.

[QUESTION] With the default parameters of *MaxPooling2D()*, what do you expect on the shape of the output tensor?

[ANSWER] With those default parameters, after a *MaxPooling2D()*, the output tensor will have the integer part of the half of both the height and width of the input tensor.

[QUESTION] Try to build a first CNN model with this architecture:

```
x --> Conv2D (32 filters) ==> MaxPooling ==> Conv2D (64 filters) ==> MaxPooling ==> Flatten ==> Dense (4) --> y
```

#Your code here

```
# (1) DEFINE THE ARCHITECTURE OF MY MODEL
```

```
#first, I define all the layers and the way they are connected
```

```
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer
```

```
x = tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu')(inputs)
```

```
x = tf.keras.layers.MaxPooling2D()(x)
```

```
x = tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu')(x)
```

```
x = tf.keras.layers.MaxPooling2D()(x)
```

```
x = tf.keras.layers.Flatten()(x) #cf. question below...
```

```
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
```

```
#Then, I define my model with the input layer, the output layer and a name
```

```
my_cnn_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_cnn_model")
```

```
#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
```

```
my_cnn_model.summary()
```

```
# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
```

```
...
```

```
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
```

```
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
```

```
(3) Set the final performance metric to evaluate the model
```

```
...
```

```
my_cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
```

```
...
```

```
(1) Set the number of epochs
```

```
(2) Set the size of the (mini)batch
```

```
(3) Set the training dataset ==> here, X_train with Y_train
```

```

(4) Set the validation dataset (X_val, Y_val)
...

nb_epochs=7
batch_size=100
training_history = my_cnn_model.fit(X_train,Y_train,
                                   validation_data=(X_val, Y_val),
                                   epochs=nb_epochs,
                                   batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_cnn_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_cnn_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

```

Model: "my_cnn_model"

Layer (type)	Output Shape	Param #
input_layer_12 (InputLayer)	(None, 32, 32, 3)	0
conv2d_4 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_5 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_10 (Flatten)	(None, 2304)	0
dense_44 (Dense)	(None, 4)	9,220

Total params: 28,612 (111.77 KB)

Trainable params: 28,612 (111.77 KB)

Non-trainable params: 0 (0.00 B)

```

Epoch 1/7
168/168 — 3s 8ms/step - accuracy: 0.3539 - loss: 10.3783 - val_accuracy: 0.5342 - val_loss: 1.2216
Epoch 2/7
168/168 — 1s 4ms/step - accuracy: 0.5634 - loss: 1.1031 - val_accuracy: 0.6064 - val_loss: 0.9867
Epoch 3/7
168/168 — 1s 4ms/step - accuracy: 0.6311 - loss: 0.9360 - val_accuracy: 0.6553 - val_loss: 0.8853
Epoch 4/7
168/168 — 1s 3ms/step - accuracy: 0.6785 - loss: 0.8200 - val_accuracy: 0.6583 - val_loss: 0.8801
Epoch 5/7
168/168 — 1s 4ms/step - accuracy: 0.7027 - loss: 0.7615 - val_accuracy: 0.6781 - val_loss: 0.8426
Epoch 6/7
168/168 — 1s 3ms/step - accuracy: 0.7157 - loss: 0.7308 - val_accuracy: 0.6881 - val_loss: 0.8145
Epoch 7/7
168/168 — 1s 3ms/step - accuracy: 0.7418 - loss: 0.6742 - val_accuracy: 0.6944 - val_loss: 0.8179
168/168 — 0s 2ms/step - accuracy: 0.7625 - loss: 0.6317
36/36 — 0s 2ms/step - accuracy: 0.7043 - loss: 0.8290
Performance on the TRAIN set, ACCURACY= 0.7643452286720276
Performance on the TEST set, ACCURACY= 0.6944444179534012

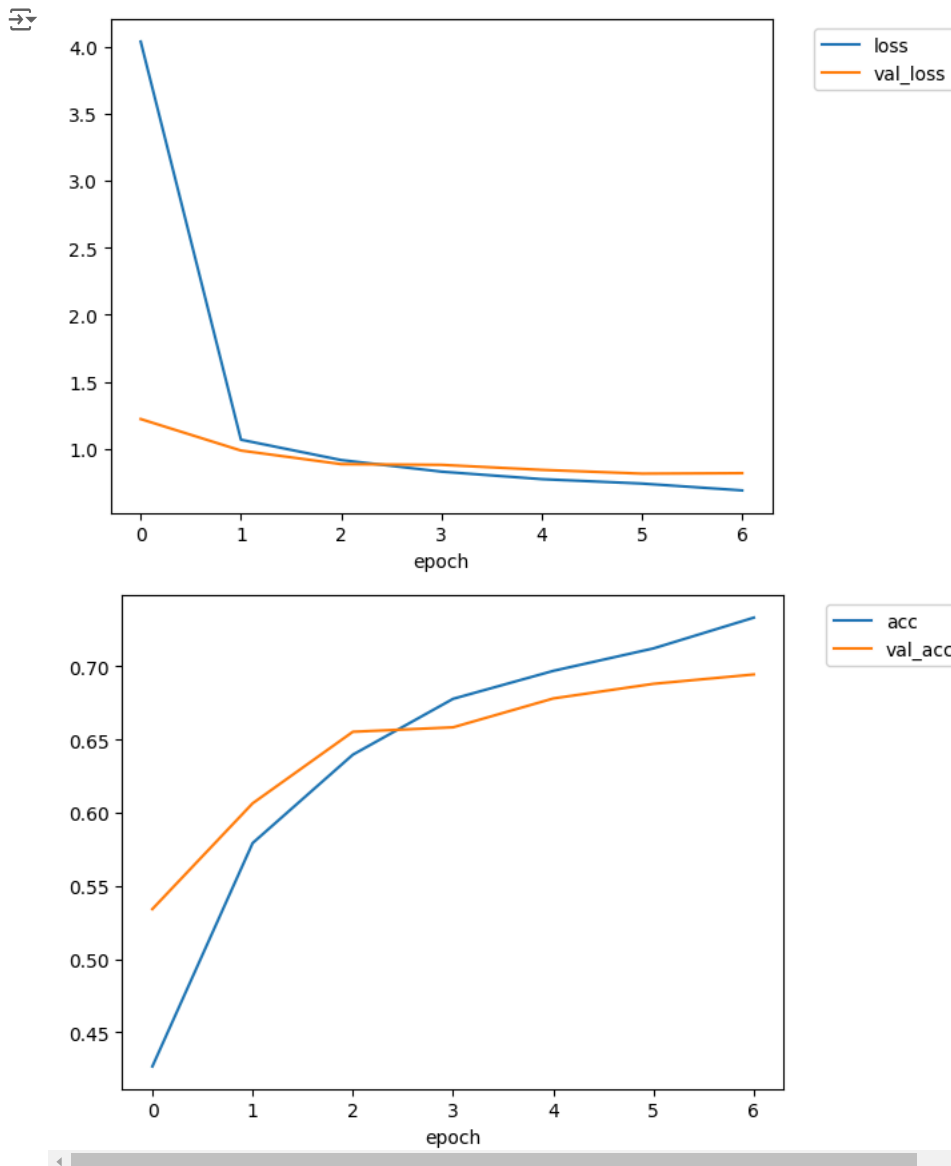
```

[QUESTION] Display the training curves from your CNN model.

```

#Your code here
display_training_curves(training_history)

```



✓ Improvements and overfitting [6 pts]

[QUESTION] Use, the code of your last model and adapt it so that your model **OVERFITS**.

We need to see the overfitting issue on the training curve!

#Your code here

```
# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer

x = tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu')(inputs)
x = tf.keras.layers.MaxPooling2D()(x)
x = tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu')(x)
x = tf.keras.layers.MaxPooling2D()(x)

x = tf.keras.layers.Flatten()(x) #cf. question below...

outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_cnn_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_cnn_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_cnn_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''
```

```
my_cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=20
batch_size=100
training_history = my_cnn_model.fit(X_train,Y_train,
                                    validation_data=(X_val, Y_val),
                                    epochs=nb_epochs,
                                    batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_cnn_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_cnn_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

display_training_curves(training_history)
```

Model: "my_cnn_model"

Layer (type)	Output Shape	Param #
input_layer_13 (InputLayer)	(None, 32, 32, 3)	0
conv2d_6 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_6 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_7 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_11 (Flatten)	(None, 2304)	0
dense_45 (Dense)	(None, 4)	9,220

Total params: 28,612 (111.77 KB)

Trainable params: 28,612 (111.77 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/20

168/168 ————— 3s 8ms/step - accuracy: 0.3548 - loss: 8.2875 - val_accuracy: 0.5075 - val_loss: 1.2061

Epoch 2/20

168/168 ————— 1s 4ms/step - accuracy: 0.5358 - loss: 1.1131 - val_accuracy: 0.6061 - val_loss: 0.9752

Epoch 3/20

168/168 ————— 1s 3ms/step - accuracy: 0.6153 - loss: 0.9523 - val_accuracy: 0.6100 - val_loss: 0.9490

Epoch 4/20

168/168 ————— 1s 4ms/step - accuracy: 0.6614 - loss: 0.8515 - val_accuracy: 0.6353 - val_loss: 0.9043

Epoch 5/20

168/168 ————— 1s 4ms/step - accuracy: 0.6810 - loss: 0.8010 - val_accuracy: 0.6514 - val_loss: 0.9022

Epoch 6/20

168/168 ————— 1s 4ms/step - accuracy: 0.7093 - loss: 0.7509 - val_accuracy: 0.6706 - val_loss: 0.8494

Epoch 7/20

168/168 ————— 2s 5ms/step - accuracy: 0.7331 - loss: 0.6952 - val_accuracy: 0.6778 - val_loss: 0.8303

Epoch 8/20

168/168 ————— 1s 6ms/step - accuracy: 0.7440 - loss: 0.6537 - val_accuracy: 0.6847 - val_loss: 0.8345

Epoch 9/20

168/168 ————— 1s 5ms/step - accuracy: 0.7665 - loss: 0.6154 - val_accuracy: 0.6967 - val_loss: 0.8152

Epoch 10/20

168/168 ————— 1s 6ms/step - accuracy: 0.7847 - loss: 0.5678 - val_accuracy: 0.6869 - val_loss: 0.8163

Epoch 11/20

168/168 ————— 1s 4ms/step - accuracy: 0.7852 - loss: 0.5597 - val_accuracy: 0.7000 - val_loss: 0.8274

Epoch 12/20

168/168 ————— 1s 4ms/step - accuracy: 0.7922 - loss: 0.5427 - val_accuracy: 0.7139 - val_loss: 0.8000

Epoch 13/20

168/168 ————— 1s 4ms/step - accuracy: 0.8181 - loss: 0.4900 - val_accuracy: 0.6842 - val_loss: 0.9031

Epoch 14/20

168/168 ————— 1s 4ms/step - accuracy: 0.8181 - loss: 0.4761 - val_accuracy: 0.6839 - val_loss: 0.9418

Epoch 15/20

168/168 ————— 1s 4ms/step - accuracy: 0.8144 - loss: 0.4786 - val_accuracy: 0.7083 - val_loss: 0.8744

Epoch 16/20

168/168 ————— 1s 3ms/step - accuracy: 0.8262 - loss: 0.4501 - val_accuracy: 0.6956 - val_loss: 0.9117

Epoch 17/20

168/168 ————— 1s 4ms/step - accuracy: 0.8361 - loss: 0.4351 - val_accuracy: 0.6931 - val_loss: 0.9097

Epoch 18/20

168/168 ————— 1s 4ms/step - accuracy: 0.8418 - loss: 0.4140 - val_accuracy: 0.7075 - val_loss: 0.8958

Epoch 19/20

168/168 ————— 1s 4ms/step - accuracy: 0.8525 - loss: 0.3881 - val_accuracy: 0.7011 - val_loss: 0.9700

Epoch 20/20


168/168 ————— 1s 4ms/step - accuracy: 0.8542 - loss: 0.3816 - val_accuracy: 0.7031 - val_loss: 1.0137


168/168 ————— 0s 2ms/step - accuracy: 0.8662 - loss: 0.3551

36/36 ————— 0s 2ms/step - accuracy: 0.7129 - loss: 0.9700

Performance on the TRAIN set, ACCURACY= 0.8677380681037903

Performance on the TEST set, ACCURACY= 0.7036111354827881





[QUESTION] Fix your overfitting issue with a technique of your choice.

#Your code here

#To fight overfitting with CNN, we can use the exact same techniques of dropping out, regularization and batch normalization

#this time let's do data augmentation

from tensorflow.keras.preprocessing.image import ImageDataGenerator

#here we can see that data augmentation is by far the best method to fight overfitting, and at the same time it improves the performance

```
datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(X_train)
```

(1) DEFINE THE ARCHITECTURE OF MY MODEL

#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer

```
x = tf.keras.layers.Conv2D(32,kernel_size=3,activation='relu')(inputs)
x = tf.keras.layers.MaxPooling2D()(x)
x = tf.keras.layers.Conv2D(64,kernel_size=3,activation='relu')(x)
x = tf.keras.layers.MaxPooling2D()(x)
```

x = tf.keras.layers.Flatten()(x) #cf. question below...

outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer

#Then, I define my model with the input layer, the output layer and a name

my_cnn_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_cnn_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS

my_cnn_model.summary()

(2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:

```
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''
```

my_cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

(3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD

```
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''
```

nb_epochs=20

batch_size=100

training_history = my_cnn_model.fit(datagen.flow(X_train, Y_train, batch_size=batch_size),


```
validation_data=(X_val, Y_val),
epochs=nb_epochs)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_cnn_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_cnn_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

display_training_curves(training_history)
```

Model: "my_cnn_model"

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 32, 32, 3)	0
conv2d_2 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_3 (Flatten)	(None, 2304)	0
dense_5 (Dense)	(None, 4)	9,220

Total params: 28,612 (111.77 KB)

Trainable params: 28,612 (111.77 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/20

168/168 ————— 14s 71ms/step - accuracy: 0.3367 - loss: 7.4771 - val_accuracy: 0.4758 - val_loss: 1.2342

Epoch 2/20

168/168 ————— 19s 67ms/step - accuracy: 0.4642 - loss: 1.2131 - val_accuracy: 0.5058 - val_loss: 1.1396

Epoch 3/20

168/168 ————— 12s 71ms/step - accuracy: 0.5416 - loss: 1.0850 - val_accuracy: 0.6175 - val_loss: 0.9666

Epoch 4/20

168/168 ————— 13s 73ms/step - accuracy: 0.5852 - loss: 1.0016 - val_accuracy: 0.6169 - val_loss: 0.9661

Epoch 5/20

168/168 ————— 12s 71ms/step - accuracy: 0.6103 - loss: 0.9665 - val_accuracy: 0.6353 - val_loss: 0.8834

Epoch 6/20

168/168 ————— 21s 73ms/step - accuracy: 0.6282 - loss: 0.9155 - val_accuracy: 0.6744 - val_loss: 0.8637

Epoch 7/20

168/168 ————— 13s 74ms/step - accuracy: 0.6332 - loss: 0.9120 - val_accuracy: 0.6678 - val_loss: 0.8380

Epoch 8/20

168/168 ————— 16s 96ms/step - accuracy: 0.6359 - loss: 0.9035 - val_accuracy: 0.6819 - val_loss: 0.8045

Epoch 9/20

168/168 ————— 13s 77ms/step - accuracy: 0.6563 - loss: 0.8720 - val_accuracy: 0.7028 - val_loss: 0.7556

Epoch 10/20

168/168 ————— 22s 87ms/step - accuracy: 0.6760 - loss: 0.8264 - val_accuracy: 0.7008 - val_loss: 0.7636

Epoch 11/20

168/168 ————— 14s 80ms/step - accuracy: 0.6801 - loss: 0.8165 - val_accuracy: 0.6889 - val_loss: 0.8028

Epoch 12/20

168/168 ————— 12s 67ms/step - accuracy: 0.6791 - loss: 0.8250 - val_accuracy: 0.7108 - val_loss: 0.7325

Epoch 13/20

168/168 ————— 23s 81ms/step - accuracy: 0.6987 - loss: 0.7810 - val_accuracy: 0.7103 - val_loss: 0.7322

Epoch 14/20

168/168 ————— 13s 75ms/step - accuracy: 0.6903 - loss: 0.7963 - val_accuracy: 0.6808 - val_loss: 0.8543

Epoch 15/20

168/168 ————— 13s 73ms/step - accuracy: 0.7001 - loss: 0.7627 - val_accuracy: 0.7156 - val_loss: 0.7209

Epoch 16/20

168/168 ————— 12s 71ms/step - accuracy: 0.6853 - loss: 0.7938 - val_accuracy: 0.7219 - val_loss: 0.7040

Epoch 17/20

168/168 ————— 12s 70ms/step - accuracy: 0.7038 - loss: 0.7621 - val_accuracy: 0.7064 - val_loss: 0.7565

Epoch 18/20

168/168 ————— 22s 79ms/step - accuracy: 0.6979 - loss: 0.7762 - val_accuracy: 0.7342 - val_loss: 0.6972

Epoch 19/20

168/168 ————— 13s 75ms/step - accuracy: 0.7180 - loss: 0.7384 - val_accuracy: 0.7108 - val_loss: 0.7357

Epoch 20/20


168/168 ————— 14s 84ms/step - accuracy: 0.7100 - loss: 0.7461 - val_accuracy: 0.7206 - val_loss: 0.7230


168/168 ————— 1s 3ms/step - accuracy: 0.7230 - loss: 0.7153

36/36 ————— 0s 3ms/step - accuracy: 0.7130 - loss: 0.7350

Performance on the TRAIN set, ACCURACY= 0.7289285659790039

Performance on the TEST set, ACCURACY= 0.7149999737739563





[QUESTION] Regarding these experiences, compare MLP vs. CNN

[ANSWER] In the CIFAR-4 classification task, **CNNs consistently outperform MLPs in terms of accuracy and parameter efficiency**. A deep MLP with 4.5M parameters achieved a test accuracy of 58% after 60 epochs, with significant overfitting, as the training accuracy reached 77%. In contrast, a simple CNN architecture with only 28K parameters achieved 71% test accuracy using data augmentation and ~70% without augmentation after 20 epochs. The training-test accuracy gap for CNNs was narrower (72% vs. 71%) compared to MLPs (77% vs. 58%), indicating better generalization. Data augmentation further improved CNN performance, reducing overfitting and increasing robustness, albeit with significantly longer training times (up to 18x slower per epoch). Thus, CNNs demonstrated superior performance, higher parameter efficiency, and better generalization on CIFAR-4 compared to MLPs.

✓ BONUS

Double-cliquez (ou appuyez sur Entrée) pour modifier

#Your code here

#To fight overfitting with CNN, we can use the exact same techniques of dropping out, regularization and batch normalization

#this time let's do data augmentation
from tensorflow.keras.preprocessing.image import ImageDataGenerator

#here we can see that data augmentation is by far the best method to fight overfitting, and at the same time it improves the performance

```
datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(X_train)
```

(1) DEFINE THE ARCHITECTURE OF MY MODEL

#first, I define all the layers and the way they are connected
inputs = tf.keras.Input(shape=(32,32,3)) #my input layer

```
x = tf.keras.layers.Conv2D(32, kernel_size=3, activation='relu')(inputs)
x = tf.keras.layers.MaxPooling2D()(x)
x = tf.keras.layers.Conv2D(64, kernel_size=3, activation='relu')(x)
x = tf.keras.layers.MaxPooling2D()(x)
```

x = tf.keras.layers.Flatten()(x) #cf. question below...

```
outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_cnn_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_cnn_model")
```

```
#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_cnn_model.summary()
```

(2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:

```
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=60
batch_size=100
training_history = my_cnn_model.fit(datagen.flow(X_train, Y_train, batch_size=batch_size),
                                   validation_data=(X_val, Y_val),
                                   epochs=nb_epochs)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_cnn_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_cnn_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)
```

Model: "my_cnn_model"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 32, 32, 3)	0
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_2 (Flatten)	(None, 2304)	0
dense_4 (Dense)	(None, 4)	9,220

Total params: 28,612 (111.77 KB)

Trainable params: 28,612 (111.77 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/60

/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: Your `PyDataset` c
self._warn_if_super_not_called()

168/168 ————— 17s 79ms/step - accuracy: 0.3135 - loss: 7.2477 - val_accuracy: 0.4542 - val_loss: 1.3439

Epoch 2/60

168/168 ————— 16s 65ms/step - accuracy: 0.4473 - loss: 1.2496 - val_accuracy: 0.5197 - val_loss: 1.1446

Epoch 3/60

168/168 ————— 21s 70ms/step - accuracy: 0.5069 - loss: 1.1427 - val_accuracy: 0.5542 - val_loss: 1.1192

Epoch 4/60

168/168 ————— 20s 70ms/step - accuracy: 0.5410 - loss: 1.0860 - val_accuracy: 0.5714 - val_loss: 1.0379

Epoch 5/60

168/168 ————— 20s 69ms/step - accuracy: 0.5670 - loss: 1.0420 - val_accuracy: 0.5833 - val_loss: 1.0318

Epoch 6/60

168/168 ————— 12s 69ms/step - accuracy: 0.5789 - loss: 1.0128 - val_accuracy: 0.6103 - val_loss: 0.9751

Epoch 7/60

168/168 ————— 12s 69ms/step - accuracy: 0.6042 - loss: 0.9691 - val_accuracy: 0.6678 - val_loss: 0.8670

Epoch 8/60

168/168 ————— 22s 77ms/step - accuracy: 0.6244 - loss: 0.9343 - val_accuracy: 0.6467 - val_loss: 0.8871

Epoch 9/60

168/168 ————— 13s 72ms/step - accuracy: 0.6365 - loss: 0.9064 - val_accuracy: 0.6500 - val_loss: 0.8839

Epoch 10/60

168/168 ————— 20s 70ms/step - accuracy: 0.6504 - loss: 0.8814 - val_accuracy: 0.6600 - val_loss: 0.8826

Epoch 11/60

168/168 ————— 20s 67ms/step - accuracy: 0.6512 - loss: 0.8685 - val_accuracy: 0.6931 - val_loss: 0.8274

Epoch 12/60

168/168 ————— 21s 72ms/step - accuracy: 0.6738 - loss: 0.8271 - val_accuracy: 0.6572 - val_loss: 0.8860

Epoch 13/60

168/168 ————— 20s 69ms/step - accuracy: 0.6590 - loss: 0.8582 - val_accuracy: 0.7061 - val_loss: 0.7618

Epoch 14/60

168/168 ————— 12s 70ms/step - accuracy: 0.6788 - loss: 0.8134 - val_accuracy: 0.6778 - val_loss: 0.8451

Epoch 15/60

168/168 ————— 12s 69ms/step - accuracy: 0.6972 - loss: 0.7760 - val_accuracy: 0.6797 - val_loss: 0.7980

Epoch 16/60

168/168 ————— 20s 65ms/step - accuracy: 0.6911 - loss: 0.7926 - val_accuracy: 0.6911 - val_loss: 0.8060

Epoch 17/60

168/168 ————— 12s 66ms/step - accuracy: 0.6966 - loss: 0.7804 - val_accuracy: 0.7083 - val_loss: 0.7394

Epoch 18/60

168/168 ————— 12s 71ms/step - accuracy: 0.6972 - loss: 0.7753 - val_accuracy: 0.7275 - val_loss: 0.7265

Epoch 19/60

168/168 ————— 20s 70ms/step - accuracy: 0.6998 - loss: 0.7713 - val_accuracy: 0.7400 - val_loss: 0.6911

Epoch 20/60

168/168 ————— 11s 64ms/step - accuracy: 0.6930 - loss: 0.7775 - val_accuracy: 0.6833 - val_loss: 0.8025

Epoch 21/60

168/168 ————— 12s 66ms/step - accuracy: 0.7052 - loss: 0.7518 - val_accuracy: 0.7322 - val_loss: 0.6978

Epoch 22/60

168/168 ————— 12s 70ms/step - accuracy: 0.7137 - loss: 0.7361 - val_accuracy: 0.7414 - val_loss: 0.6628

Epoch 23/60

168/168 ————— 12s 69ms/step - accuracy: 0.7080 - loss: 0.7442 - val_accuracy: 0.6933 - val_loss: 0.8086

Epoch 24/60

168/168 ————— 20s 65ms/step - accuracy: 0.7072 - loss: 0.7450 - val_accuracy: 0.7369 - val_loss: 0.6816

Epoch 25/60

168/168 ————— 11s 62ms/step - accuracy: 0.7141 - loss: 0.7342 - val_accuracy: 0.7592 - val_loss: 0.6395

Epoch 26/60

168/168 ————— 12s 71ms/step - accuracy: 0.7213 - loss: 0.7090 - val_accuracy: 0.7417 - val_loss: 0.6538

Epoch 27/60

168/168 ————— 12s 70ms/step - accuracy: 0.7246 - loss: 0.7144 - val_accuracy: 0.7656 - val_loss: 0.6396

Epoch 28/60

168/168 ————— 20s 70ms/step - accuracy: 0.7249 - loss: 0.7150 - val_accuracy: 0.7269 - val_loss: 0.6930

Epoch 29/60

168/168 ————— 20s 70ms/step - accuracy: 0.7283 - loss: 0.7125 - val_accuracy: 0.7564 - val_loss: 0.6480

Epoch 30/60

168/168 ————— 22s 80ms/step - accuracy: 0.7238 - loss: 0.7201 - val_accuracy: 0.7406 - val_loss: 0.6879

Epoch 31/60

168/168 ————— 19s 70ms/step - accuracy: 0.7251 - loss: 0.7069 - val_accuracy: 0.7392 - val_loss: 0.6660

Epoch 32/60

168/168 ————— 20s 65ms/step - accuracy: 0.7336 - loss: 0.6938 - val_accuracy: 0.7397 - val_loss: 0.7212

Epoch 33/60

168/168 ————— 12s 67ms/step - accuracy: 0.7238 - loss: 0.7134 - val_accuracy: 0.7433 - val_loss: 0.6885

Epoch 34/60

168/168 ————— 12s 72ms/step - accuracy: 0.7350 - loss: 0.6811 - val_accuracy: 0.7336 - val_loss: 0.7330

```

Epoch 35/60
168/168 12s 70ms/step - accuracy: 0.7281 - loss: 0.6987 - val_accuracy: 0.7344 - val_loss: 0.7395
Epoch 36/60
168/168 20s 69ms/step - accuracy: 0.7367 - loss: 0.6775 - val_accuracy: 0.7181 - val_loss: 0.7447
Epoch 37/60
168/168 21s 69ms/step - accuracy: 0.7363 - loss: 0.6894 - val_accuracy: 0.7697 - val_loss: 0.6247
Epoch 38/60
168/168 12s 69ms/step - accuracy: 0.7455 - loss: 0.6640 - val_accuracy: 0.7661 - val_loss: 0.6344
Epoch 39/60
168/168 22s 78ms/step - accuracy: 0.7423 - loss: 0.6678 - val_accuracy: 0.7606 - val_loss: 0.6341
Epoch 40/60
168/168 11s 63ms/step - accuracy: 0.7535 - loss: 0.6560 - val_accuracy: 0.7472 - val_loss: 0.6731
Epoch 41/60
168/168 18s 106ms/step - accuracy: 0.7465 - loss: 0.6676 - val_accuracy: 0.7600 - val_loss: 0.6520
Epoch 42/60
168/168 13s 64ms/step - accuracy: 0.7423 - loss: 0.6751 - val_accuracy: 0.7525 - val_loss: 0.6632
Epoch 43/60
168/168 11s 61ms/step - accuracy: 0.7403 - loss: 0.6807 - val_accuracy: 0.7786 - val_loss: 0.5962
Epoch 44/60
168/168 12s 69ms/step - accuracy: 0.7549 - loss: 0.6478 - val_accuracy: 0.7511 - val_loss: 0.6787
Epoch 45/60
168/168 12s 68ms/step - accuracy: 0.7531 - loss: 0.6516 - val_accuracy: 0.7533 - val_loss: 0.6858
Epoch 46/60
168/168 20s 67ms/step - accuracy: 0.7554 - loss: 0.6409 - val_accuracy: 0.7408 - val_loss: 0.7059
Epoch 47/60
168/168 21s 70ms/step - accuracy: 0.7423 - loss: 0.6687 - val_accuracy: 0.7572 - val_loss: 0.6640
Epoch 48/60
168/168 12s 68ms/step - accuracy: 0.7542 - loss: 0.6424 - val_accuracy: 0.7653 - val_loss: 0.6113
Epoch 49/60
168/168 12s 70ms/step - accuracy: 0.7645 - loss: 0.6246 - val_accuracy: 0.7617 - val_loss: 0.6484
Epoch 50/60
168/168 12s 69ms/step - accuracy: 0.7508 - loss: 0.6543 - val_accuracy: 0.7703 - val_loss: 0.6208
Epoch 51/60
168/168 11s 61ms/step - accuracy: 0.7575 - loss: 0.6506 - val_accuracy: 0.7708 - val_loss: 0.5898
Epoch 52/60
168/168 12s 66ms/step - accuracy: 0.7579 - loss: 0.6360 - val_accuracy: 0.7794 - val_loss: 0.6029
Epoch 53/60
168/168 21s 69ms/step - accuracy: 0.7551 - loss: 0.6446 - val_accuracy: 0.7494 - val_loss: 0.6826
Epoch 54/60
168/168 27s 100ms/step - accuracy: 0.7494 - loss: 0.6531 - val_accuracy: 0.7594 - val_loss: 0.6410
Epoch 55/60
168/168 13s 73ms/step - accuracy: 0.7438 - loss: 0.6568 - val_accuracy: 0.7747 - val_loss: 0.6157
Epoch 56/60
168/168 16s 92ms/step - accuracy: 0.7693 - loss: 0.6098 - val_accuracy: 0.7581 - val_loss: 0.6375
Epoch 57/60
168/168 16s 91ms/step - accuracy: 0.7542 - loss: 0.6430 - val_accuracy: 0.7592 - val_loss: 0.6446
Epoch 58/60
168/168 13s 73ms/step - accuracy: 0.7572 - loss: 0.6354 - val_accuracy: 0.7514 - val_loss: 0.6911
Epoch 59/60
168/168 11s 60ms/step - accuracy: 0.7621 - loss: 0.6377 - val_accuracy: 0.7481 - val_loss: 0.6875
Epoch 60/60
168/168 12s 71ms/step - accuracy: 0.7512 - loss: 0.6464 - val_accuracy: 0.7789 - val_loss: 0.5801
168/168 1s 2ms/step - accuracy: 0.8020 - loss: 0.5307
36/36 0s 2ms/step - accuracy: 0.7853 - loss: 0.5974
Performance on the TRAIN set, ACCURACY= 0.8045833110809326
Performance on the TEST set, ACCURACY= 0.7827777862548828

```

```


-----
NameError                                Traceback (most recent call last)
<ipython-input-12-ed0303bc2989> in <cell line: 66>()
    64 print("Performance on the TEST set, ACCURACY=",acc_test)
    65
--> 66 display_training_curves(training_history)
    67

```

NameError: name 'display_training_curves' is not defined

Étapes suivantes : [Expliquer l'erreur](#)

display_training_curves(training_history)



[ANSWER] This CNN achieved ~78.3% test accuracy, with data augmentation and **60 epochs** eliminating completely the gap generalization for 60 epochs or less, albeit at 18x slower training time. This simple architecture (28,612 parameters) effectively captured spatial patterns, proving CNNs' strength in image classification by an outstanding performance of 78%. Which is the best result we get through all of our different MLP algorithms.

#Your code here

#Here we will do some transfer learning with the ResNet50 algorithm, a pre-trained deep convolutional neural network, for our CIFAR-4 dataset

```
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import Dense, Flatten, GlobalAveragePooling2D, Input

# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected

# Load the ResNet50 model pre-trained on ImageNet
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

# Freeze the base model (optional if you only want to train the top layers)
base_model.trainable = False

# Add custom layers for CIFAR-4 classification
inputs = base_model.input
x = base_model.output
x = GlobalAveragePooling2D()(x) # Pooling to reduce dimensions
x = Dense(256, activation='relu')(x) # Fully connected layer
```

```

x = Dense(128, activation='relu')(x) # Fully connected layer

outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_res50_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_res50_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_res50_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_res50_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''


nb_epochs=15
batch_size=100
training_history = my_res50_model.fit(datagen.flow(X_train, Y_train, batch_size=batch_size),
                                     validation_data=(X_val, Y_val),
                                     epochs=nb_epochs,
                                     batch_size=batch_size)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_res50_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_res50_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

display_training_curves(training_history)

```


 Model: "my_res50_model"

Layer (type)	Output Shape	Param #	Connected to
input_layer_14 (InputLayer)	(None, 32, 32, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 38, 38, 3)	0	input_layer_14[0][0]
conv1_conv (Conv2D)	(None, 16, 16, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 16, 16, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 18, 18, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 8, 8, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 8, 8, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block1_1_conv[0...
conv2_block1_1_relu (Activation)	(None, 8, 8, 64)	0	conv2_block1_1_bn[0][...
conv2_block1_2_conv (Conv2D)	(None, 8, 8, 64)	36,928	conv2_block1_1_relu[0...
conv2_block1_2_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block1_2_conv[0...
conv2_block1_2_relu (Activation)	(None, 8, 8, 64)	0	conv2_block1_2_bn[0][...
conv2_block1_0_conv (Conv2D)	(None, 8, 8, 256)	16,640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 8, 8, 256)	16,640	conv2_block1_2_relu[0...
conv2_block1_0_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block1_0_conv[0...
conv2_block1_3_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block1_3_conv[0...
conv2_block1_add (Add)	(None, 8, 8, 256)	0	conv2_block1_0_bn[0][...] conv2_block1_3_bn[0][...]
conv2_block1_out (Activation)	(None, 8, 8, 256)	0	conv2_block1_add[0][0]
conv2_block2_1_conv (Conv2D)	(None, 8, 8, 64)	16,448	conv2_block1_out[0][0]
conv2_block2_1_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block2_1_conv[0...
conv2_block2_1_relu (Activation)	(None, 8, 8, 64)	0	conv2_block2_1_bn[0][...]
conv2_block2_2_conv (Conv2D)	(None, 8, 8, 64)	36,928	conv2_block2_1_relu[0...
conv2_block2_2_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block2_2_conv[0...
conv2_block2_2_relu (Activation)	(None, 8, 8, 64)	0	conv2_block2_2_bn[0][...]
conv2_block2_3_conv (Conv2D)	(None, 8, 8, 256)	16,640	conv2_block2_2_relu[0...
conv2_block2_3_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block2_3_conv[0...
conv2_block2_add (Add)	(None, 8, 8, 256)	0	conv2_block1_out[0][0]... conv2_block2_3_bn[0][...]
conv2_block2_out (Activation)	(None, 8, 8, 256)	0	conv2_block2_add[0][0]
conv2_block3_1_conv (Conv2D)	(None, 8, 8, 64)	16,448	conv2_block2_out[0][0]
conv2_block3_1_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block3_1_conv[0...

conv2_block3_1_relu (Activation)	(None, 8, 8, 64)	0	conv2_block3_1_bn[0][...
conv2_block3_2_conv (Conv2D)	(None, 8, 8, 64)	36,928	conv2_block3_1_relu[0][...]
conv2_block3_2_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block3_2_conv[0][...]
conv2_block3_2_relu (Activation)	(None, 8, 8, 64)	0	conv2_block3_2_bn[0][...]
conv2_block3_3_conv (Conv2D)	(None, 8, 8, 256)	16,640	conv2_block3_2_relu[0][...]
conv2_block3_3_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block3_3_conv[0][...]
conv2_block3_add (Add)	(None, 8, 8, 256)	0	conv2_block2_out[0][0]... conv2_block3_3_bn[0][...]
conv2_block3_out (Activation)	(None, 8, 8, 256)	0	conv2_block3_add[0][0]
conv3_block1_1_conv (Conv2D)	(None, 4, 4, 128)	32,896	conv2_block3_out[0][0]
conv3_block1_1_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block1_1_conv[0][...]
conv3_block1_1_relu (Activation)	(None, 4, 4, 128)	0	conv3_block1_1_bn[0][...]
conv3_block1_2_conv (Conv2D)	(None, 4, 4, 128)	147,584	conv3_block1_1_relu[0][...]
conv3_block1_2_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block1_2_conv[0][...]
conv3_block1_2_relu (Activation)	(None, 4, 4, 128)	0	conv3_block1_2_bn[0][...]
conv3_block1_0_conv (Conv2D)	(None, 4, 4, 512)	131,584	conv2_block3_out[0][0]
conv3_block1_3_conv (Conv2D)	(None, 4, 4, 512)	66,048	conv3_block1_2_relu[0][...]
conv3_block1_0_bn (BatchNormalization)	(None, 4, 4, 512)	2,048	conv3_block1_0_conv[0][...]
conv3_block1_3_bn (BatchNormalization)	(None, 4, 4, 512)	2,048	conv3_block1_3_conv[0][...]
conv3_block1_add (Add)	(None, 4, 4, 512)	0	conv3_block1_0_bn[0][...] conv3_block1_3_bn[0][...]
conv3_block1_out (Activation)	(None, 4, 4, 512)	0	conv3_block1_add[0][0]
conv3_block2_1_conv (Conv2D)	(None, 4, 4, 128)	65,664	conv3_block1_out[0][0]
conv3_block2_1_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block2_1_conv[0][...]
conv3_block2_1_relu (Activation)	(None, 4, 4, 128)	0	conv3_block2_1_bn[0][...]
conv3_block2_2_conv (Conv2D)	(None, 4, 4, 128)	147,584	conv3_block2_1_relu[0][...]
conv3_block2_2_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block2_2_conv[0][...]
conv3_block2_2_relu (Activation)	(None, 4, 4, 128)	0	conv3_block2_2_bn[0][...]
conv3_block2_3_conv (Conv2D)	(None, 4, 4, 512)	66,048	conv3_block2_2_relu[0][...]
conv3_block2_3_bn (BatchNormalization)	(None, 4, 4, 512)	2,048	conv3_block2_3_conv[0][...]
conv3_block2_add (Add)	(None, 4, 4, 512)	0	conv3_block1_out[0][0]... conv3_block2_3_bn[0][...]
conv3_block2_out (Activation)	(None, 4, 4, 512)	0	conv3_block2_add[0][0]
conv3_block3_1_conv (Conv2D)	(None, 4, 4, 128)	65,664	conv3_block2_out[0][0]

conv3_block3_1_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block3_1_conv[0...
conv3_block3_1_relu (Activation)	(None, 4, 4, 128)	0	conv3_block3_1_bn[0][...
conv3_block3_2_conv (Conv2D)	(None, 4, 4, 128)	147,584	conv3_block3_1_relu[0...
conv3_block3_2_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block3_2_conv[0...
conv3_block3_2_relu (Activation)	(None, 4, 4, 128)	0	conv3_block3_2_bn[0][...
conv3_block3_3_conv (Conv2D)	(None, 4, 4, 512)	66,048	conv3_block3_2_relu[0...
conv3_block3_3_bn (BatchNormalization)	(None, 4, 4, 512)	2,048	conv3_block3_3_conv[0...
conv3_block3_add (Add)	(None, 4, 4, 512)	0	conv3_block2_out[0][0... conv3_block3_3_bn[0][...
conv3_block3_out (Activation)	(None, 4, 4, 512)	0	conv3_block3_add[0][0]
conv3_block4_1_conv (Conv2D)	(None, 4, 4, 128)	65,664	conv3_block3_out[0][0]
conv3_block4_1_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block4_1_conv[0...
conv3_block4_1_relu (Activation)	(None, 4, 4, 128)	0	conv3_block4_1_bn[0][...
conv3_block4_2_conv (Conv2D)	(None, 4, 4, 128)	147,584	conv3_block4_1_relu[0...
conv3_block4_2_bn (BatchNormalization)	(None, 4, 4, 128)	512	conv3_block4_2_conv[0...
conv3_block4_2_relu (Activation)	(None, 4, 4, 128)	0	conv3_block4_2_bn[0][...
conv3_block4_3_conv (Conv2D)	(None, 4, 4, 512)	66,048	conv3_block4_2_relu[0...
conv3_block4_3_bn (BatchNormalization)	(None, 4, 4, 512)	2,048	conv3_block4_3_conv[0...
conv3_block4_add (Add)	(None, 4, 4, 512)	0	conv3_block3_out[0][0... conv3_block4_3_bn[0][...
conv3_block4_out (Activation)	(None, 4, 4, 512)	0	conv3_block4_add[0][0]
conv4_block1_1_conv (Conv2D)	(None, 2, 2, 256)	131,328	conv3_block4_out[0][0]
conv4_block1_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block1_1_conv[0...
conv4_block1_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block1_1_bn[0][...
conv4_block1_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block1_1_relu[0...
conv4_block1_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block1_2_conv[0...
conv4_block1_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block1_2_bn[0][...
conv4_block1_0_conv (Conv2D)	(None, 2, 2, 1024)	525,312	conv3_block4_out[0][0]
conv4_block1_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block1_2_relu[0...
conv4_block1_0_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block1_0_conv[0...
conv4_block1_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block1_3_conv[0...
conv4_block1_add (Add)	(None, 2, 2, 1024)	0	conv4_block1_0_bn[0][... conv4_block1_3_bn[0][...
conv4_block1_out (Activation)	(None, 2, 2, 1024)	0	conv4_block1_add[0][0]

conv4_block2_1_conv (Conv2D)	(None, 2, 2, 256)	262,400	conv4_block1_out[0][0]
conv4_block2_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block2_1_conv[0...
conv4_block2_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block2_1_bn[0][...
conv4_block2_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block2_1_relu[0...
conv4_block2_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block2_2_conv[0...
conv4_block2_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block2_2_bn[0][...
conv4_block2_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block2_2_relu[0...
conv4_block2_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block2_3_conv[0...
conv4_block2_add (Add)	(None, 2, 2, 1024)	0	conv4_block1_out[0][0... conv4_block2_3_bn[0][...
conv4_block2_out (Activation)	(None, 2, 2, 1024)	0	conv4_block2_add[0][0]
conv4_block3_1_conv (Conv2D)	(None, 2, 2, 256)	262,400	conv4_block2_out[0][0]
conv4_block3_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block3_1_conv[0...
conv4_block3_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block3_1_bn[0][...
conv4_block3_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block3_1_relu[0...
conv4_block3_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block3_2_conv[0...
conv4_block3_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block3_2_bn[0][...
conv4_block3_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block3_2_relu[0...
conv4_block3_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block3_3_conv[0...
conv4_block3_add (Add)	(None, 2, 2, 1024)	0	conv4_block2_out[0][0... conv4_block3_3_bn[0][...
conv4_block3_out (Activation)	(None, 2, 2, 1024)	0	conv4_block3_add[0][0]
conv4_block4_1_conv (Conv2D)	(None, 2, 2, 256)	262,400	conv4_block3_out[0][0]
conv4_block4_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block4_1_conv[0...
conv4_block4_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block4_1_bn[0][...
conv4_block4_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block4_1_relu[0...
conv4_block4_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block4_2_conv[0...
conv4_block4_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block4_2_bn[0][...
conv4_block4_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block4_2_relu[0...
conv4_block4_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block4_3_conv[0...
conv4_block4_add (Add)	(None, 2, 2, 1024)	0	conv4_block3_out[0][0... conv4_block4_3_bn[0][...
conv4_block4_out (Activation)	(None, 2, 2, 1024)	0	conv4_block4_add[0][0]
conv4_block5_1_conv	(None, 2, 2, 256)	262,400	conv4_block4_out[0][0]

(Conv2D)			
conv4_block5_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block5_1_conv[0...]
conv4_block5_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block5_1_bn[0][...]
conv4_block5_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block5_1_relu[0...]
conv4_block5_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block5_2_conv[0...]
conv4_block5_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block5_2_bn[0][...]
conv4_block5_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block5_2_relu[0...]
conv4_block5_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block5_3_conv[0...]
conv4_block5_add (Add)	(None, 2, 2, 1024)	0	conv4_block4_out[0][0...] conv4_block5_3_bn[0][...]
conv4_block5_out (Activation)	(None, 2, 2, 1024)	0	conv4_block5_add[0][0]
conv4_block6_1_conv (Conv2D)	(None, 2, 2, 256)	262,400	conv4_block5_out[0][0]
conv4_block6_1_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block6_1_conv[0...]
conv4_block6_1_relu (Activation)	(None, 2, 2, 256)	0	conv4_block6_1_bn[0][...]
conv4_block6_2_conv (Conv2D)	(None, 2, 2, 256)	590,080	conv4_block6_1_relu[0...]
conv4_block6_2_bn (BatchNormalization)	(None, 2, 2, 256)	1,024	conv4_block6_2_conv[0...]
conv4_block6_2_relu (Activation)	(None, 2, 2, 256)	0	conv4_block6_2_bn[0][...]
conv4_block6_3_conv (Conv2D)	(None, 2, 2, 1024)	263,168	conv4_block6_2_relu[0...]
conv4_block6_3_bn (BatchNormalization)	(None, 2, 2, 1024)	4,096	conv4_block6_3_conv[0...]
conv4_block6_add (Add)	(None, 2, 2, 1024)	0	conv4_block5_out[0][0...] conv4_block6_3_bn[0][...]
conv4_block6_out (Activation)	(None, 2, 2, 1024)	0	conv4_block6_add[0][0]
conv5_block1_1_conv (Conv2D)	(None, 1, 1, 512)	524,800	conv4_block6_out[0][0]
conv5_block1_1_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block1_1_conv[0...]
conv5_block1_1_relu (Activation)	(None, 1, 1, 512)	0	conv5_block1_1_bn[0][...]
conv5_block1_2_conv (Conv2D)	(None, 1, 1, 512)	2,359,808	conv5_block1_1_relu[0...]
conv5_block1_2_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block1_2_conv[0...]
conv5_block1_2_relu (Activation)	(None, 1, 1, 512)	0	conv5_block1_2_bn[0][...]
conv5_block1_0_conv (Conv2D)	(None, 1, 1, 2048)	2,099,200	conv4_block6_out[0][0]
conv5_block1_3_conv (Conv2D)	(None, 1, 1, 2048)	1,050,624	conv5_block1_2_relu[0...]
conv5_block1_0_bn (BatchNormalization)	(None, 1, 1, 2048)	8,192	conv5_block1_0_conv[0...]
conv5_block1_3_bn (BatchNormalization)	(None, 1, 1, 2048)	8,192	conv5_block1_3_conv[0...]
conv5_block1_add (Add)	(None, 1, 1, 2048)	0	conv5_block1_0_bn[0][...] conv5_block1_3_bn[0][...]
conv5_block1_out	(None, 1, 1, 2048)	0	conv5_block1_add[0][0]

conv5_block1_out (Activation)	(None, 1, 1, 512)	0	conv5_block1_out[0][0]
conv5_block2_1_conv (Conv2D)	(None, 1, 1, 512)	1,049,088	conv5_block1_out[0][0]
conv5_block2_1_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block2_1_conv[0...]
conv5_block2_1_relu (Activation)	(None, 1, 1, 512)	0	conv5_block2_1_bn[0][...]
conv5_block2_2_conv (Conv2D)	(None, 1, 1, 512)	2,359,808	conv5_block2_1_relu[0...]
conv5_block2_2_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block2_2_conv[0...]
conv5_block2_2_relu (Activation)	(None, 1, 1, 512)	0	conv5_block2_2_bn[0][...]
conv5_block2_3_conv (Conv2D)	(None, 1, 1, 2048)	1,050,624	conv5_block2_2_relu[0...]
conv5_block2_3_bn (BatchNormalization)	(None, 1, 1, 2048)	8,192	conv5_block2_3_conv[0...]
conv5_block2_add (Add)	(None, 1, 1, 2048)	0	conv5_block1_out[0][0...] conv5_block2_3_bn[0][...]
conv5_block2_out (Activation)	(None, 1, 1, 2048)	0	conv5_block2_add[0][0]
conv5_block3_1_conv (Conv2D)	(None, 1, 1, 512)	1,049,088	conv5_block2_out[0][0]
conv5_block3_1_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block3_1_conv[0...]
conv5_block3_1_relu (Activation)	(None, 1, 1, 512)	0	conv5_block3_1_bn[0][...]
conv5_block3_2_conv (Conv2D)	(None, 1, 1, 512)	2,359,808	conv5_block3_1_relu[0...]
conv5_block3_2_bn (BatchNormalization)	(None, 1, 1, 512)	2,048	conv5_block3_2_conv[0...]
conv5_block3_2_relu (Activation)	(None, 1, 1, 512)	0	conv5_block3_2_bn[0][...]
conv5_block3_3_conv (Conv2D)	(None, 1, 1, 2048)	1,050,624	conv5_block3_2_relu[0...]
conv5_block3_3_bn (BatchNormalization)	(None, 1, 1, 2048)	8,192	conv5_block3_3_conv[0...]
conv5_block3_add (Add)	(None, 1, 1, 2048)	0	conv5_block2_out[0][0...] conv5_block3_3_bn[0][...]
conv5_block3_out (Activation)	(None, 1, 1, 2048)	0	conv5_block3_add[0][0]
global_average_pooling2d... (GlobalAveragePooling2D)	(None, 2048)	0	conv5_block3_out[0][0]
dense_46 (Dense)	(None, 256)	524,544	global_average_poolin...
dense_47 (Dense)	(None, 128)	32,896	dense_46[0][0]
dense_48 (Dense)	(None, 4)	516	dense_47[0][0]

Total params: 24,145,668 (92.11 MB)

Trainable params: 557,956 (2.13 MB)

Non-trainable params: 23,587,712 (89.98 MB)

Epoch 1/15

168/168 ————— 28s 116ms/step - accuracy: 0.5658 - loss: 1.3415 - val_accuracy: 0.7167 - val_loss: 0.7305

Epoch 2/15

168/168 ————— 34s 91ms/step - accuracy: 0.6993 - loss: 0.7665 - val_accuracy: 0.7158 - val_loss: 0.7135

Epoch 3/15

168/168 ————— 20s 89ms/step - accuracy: 0.7225 - loss: 0.7097 - val_accuracy: 0.7386 - val_loss: 0.6840

Epoch 4/15

168/168 ————— 22s 101ms/step - accuracy: 0.7313 - loss: 0.6856 - val_accuracy: 0.7361 - val_loss: 0.6983

Epoch 5/15

168/168 ————— 19s 89ms/step - accuracy: 0.7414 - loss: 0.6591 - val_accuracy: 0.7511 - val_loss: 0.6563

Epoch 6/15

168/168 ————— 13s 76ms/step - accuracy: 0.7456 - loss: 0.6487 - val_accuracy: 0.7256 - val_loss: 0.7020

Epoch 7/15

168/168 ————— 13s 76ms/step - accuracy: 0.7505 - loss: 0.6411 - val_accuracy: 0.7556 - val_loss: 0.6459

Epoch 8/15

168/168 ————— 13s 76ms/step - accuracy: 0.7587 - loss: 0.6079 - val_accuracy: 0.7572 - val_loss: 0.6435



Epoch 9/15

168/168 ————— 14s 79ms/step - accuracy: 0.7613 - loss: 0.6142 - val_accuracy: 0.7528 - val_loss: 0.6463

```
Epoch 10/15
168/168 20s 81ms/step - accuracy: 0.7712 - loss: 0.5843 - val_accuracy: 0.7458 - val_loss: 0.6633
Epoch 11/15
168/168 21s 81ms/step - accuracy: 0.7692 - loss: 0.5860 - val_accuracy: 0.7378 - val_loss: 0.7103
Epoch 12/15
168/168 16s 89ms/step - accuracy: 0.7696 - loss: 0.5875 - val_accuracy: 0.7522 - val_loss: 0.6585
Epoch 13/15
168/168 14s 77ms/step - accuracy: 0.7782 - loss: 0.5660 - val_accuracy: 0.7564 - val_loss: 0.6603
Epoch 14/15
168/168 21s 80ms/step - accuracy: 0.7859 - loss: 0.5476 - val_accuracy: 0.7425 - val_loss: 0.7004
Epoch 15/15
168/168 16s 91ms/step - accuracy: 0.7862 - loss: 0.5481 - val_accuracy: 0.7458 - val_loss: 0.6728
168/168 2s 14ms/step - accuracy: 0.8020 - loss: 0.5024
36/36 0s 13ms/step - accuracy: 0.7318 - loss: 0.7021
```

Performance on the TRAIN set, ACCURACY= 0.8064285516738892

Performance on the TEST set, ACCURACY= 0.7313888669013977



[ANSWER] The ResNet50 model achieved a training accuracy of 80.6% and a test accuracy of 73.1% without data augmentation, highlighting its strong feature extraction capabilities due to pre-training on ImageNet. However, the 7.5% accuracy gap between training and test sets indicates some overfitting. The model's performance reflects the power of transfer learning, as the base ResNet50 layers effectively extracted features, requiring minimal fine-tuning of added dense layers. Incorporating data augmentation could further improve generalization and close the performance gap. Despite limited epochs and a frozen base model, ResNet50 efficiently adapted to the CIFAR-4 dataset, showing its robustness for smaller, specialized datasets.

#Your code here

#Here we will do some transfer learning with the ResNet50 algorithm, a pre-trained deep convolutional neural network, for our CIFAR-4 dat

```
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.layers import Dense, Flatten, GlobalAveragePooling2D, Input

datagen = ImageDataGenerator(
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True
)
datagen.fit(X_train)

# (1) DEFINE THE ARCHITECTURE OF MY MODEL
#first, I define all the layers and the way they are connected

# Load the ResNet50 model pre-trained on ImageNet
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

# Freeze the base model (optional if you only want to train the top layers)
base_model.trainable = False

# Add custom layers for CIFAR-4 classification
inputs = base_model.input
x = base_model.output
x = GlobalAveragePooling2D()(x) # Pooling to reduce dimensions
x = Dense(256, activation='relu')(x) # Fully connected layer
x = Dense(128, activation='relu')(x) # Fully connected layer

outputs = tf.keras.layers.Dense(4, activation='softmax')(x) # my output layer
#Then, I define my model with the input layer, the output layer and a name
my_res50_model = tf.keras.Model(inputs=inputs, outputs=outputs, name="my_res50_model")

#PRINT A SUMMARY OF THE ARCHITECTURE OF MY MODEL WITH THE NUMBER OF TRAINABLE PARAMETERS
my_res50_model.summary()

# (2) DEFINE THE TRAINING HYPER-PARAMETERS WITH THE "COMPILE" METHOD:
'''
(1) Set the "optimizer" [pick 'adam', 'sgd' or 'rmsprop']
(2) Set the loss [cf. lesson #3, we pick the categorical cross-entropy]
(3) Set the final performance metric to evaluate the model
'''

my_res50_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])


# (3) NOW, LET'S TRAIN ON MY DATA WITH THE "FIT" METHOD
'''
(1) Set the number of epochs
(2) Set the size of the (mini)batch
(3) Set the training dataset ==> here, X_train with Y_train
(4) Set the validation dataset (X_val, Y_val)
'''

nb_epochs=20
batch_size=100
training_history = my_res50_model.fit(datagen.flow(X_train, Y_train, batch_size=batch_size),
                                     validation_data=(X_val, Y_val),
                                     epochs=nb_epochs)

#COMPUTE THE ACCURACY ON THE TRAINING AND TEST SETS
loss_train, acc_train = my_res50_model.evaluate(X_train, Y_train, batch_size=batch_size)
loss_test, acc_test = my_res50_model.evaluate(X_test, Y_test, batch_size=batch_size)

print("Performance on the TRAIN set, ACCURACY=",acc_train)
print("Performance on the TEST set, ACCURACY=",acc_test)

display_training_curves(training_history)
```


 Model: "my_res50_model"

Layer (type)	Output Shape	Param #	Connected to
input_layer_18 (InputLayer)	(None, 32, 32, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 38, 38, 3)	0	input_layer_18[0][0]
conv1_conv (Conv2D)	(None, 16, 16, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 16, 16, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 16, 16, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 18, 18, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 8, 8, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 8, 8, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block1_1_conv[0...
conv2_block1_1_relu (Activation)	(None, 8, 8, 64)	0	conv2_block1_1_bn[0][...
conv2_block1_2_conv (Conv2D)	(None, 8, 8, 64)	36,928	conv2_block1_1_relu[0...
conv2_block1_2_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block1_2_conv[0...
conv2_block1_2_relu (Activation)	(None, 8, 8, 64)	0	conv2_block1_2_bn[0][...
conv2_block1_0_conv (Conv2D)	(None, 8, 8, 256)	16,640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 8, 8, 256)	16,640	conv2_block1_2_relu[0...
conv2_block1_0_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block1_0_conv[0...
conv2_block1_3_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block1_3_conv[0...
conv2_block1_add (Add)	(None, 8, 8, 256)	0	conv2_block1_0_bn[0][...] conv2_block1_3_bn[0][...]
conv2_block1_out (Activation)	(None, 8, 8, 256)	0	conv2_block1_add[0][0]
conv2_block2_1_conv (Conv2D)	(None, 8, 8, 64)	16,448	conv2_block1_out[0][0]
conv2_block2_1_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block2_1_conv[0...
conv2_block2_1_relu (Activation)	(None, 8, 8, 64)	0	conv2_block2_1_bn[0][...]
conv2_block2_2_conv (Conv2D)	(None, 8, 8, 64)	36,928	conv2_block2_1_relu[0...
conv2_block2_2_bn (BatchNormalization)	(None, 8, 8, 64)	256	conv2_block2_2_conv[0...
conv2_block2_2_relu (Activation)	(None, 8, 8, 64)	0	conv2_block2_2_bn[0][...]
conv2_block2_3_conv (Conv2D)	(None, 8, 8, 256)	16,640	conv2_block2_2_relu[0...
conv2_block2_3_bn (BatchNormalization)	(None, 8, 8, 256)	1,024	conv2_block2_3_conv[0...
conv2_block2_add (Add)	(None, 8, 8, 256)	0	conv2_block1_out[0][0]...

Impossible d'établir une connexion avec le service reCAPTCHA. Veuillez vérifier votre connexion Internet, puis actualiser la page pour afficher une image reCAPTCHA.