

Focus on L_k Regularization

$$x = (x_1, \dots, x_n), \text{ For } q \in \mathbb{N}^*, \|x\|_q = \sqrt[q]{\sum |x_i|^q}$$
$$\begin{cases} q=0, \|x\|_0 = \sum \mathbb{1}_{x_i \neq 0} \end{cases}$$

INDICATOR FUNCTION: $\mathbb{1}_{x_i} = \begin{cases} 0, x_i = 0 \\ 1, x_i \neq 0 \end{cases}$

I - Focus on L_0 Regularization

L_0 Regularization refers to minimizing the number of non-zero weights in a model

↳ goal: remove unnecessary parameters and make the network sparse (sparse model, interpretable, efficient)

→ effect: pruning = forces some weights to be exactly zero

⇒ Because of its discontinuous nature: Standard Gradient Descent for Regularization UNUSABLE

$$X \quad \mathcal{L} = \underbrace{\text{Loss}(W)}_{\substack{\text{total loss} \\ \text{from our NN} \\ \text{(like MSE)}}} + \underbrace{\lambda \|W\|_0}_{L_0 \text{ penalty}}$$

⇒ We need a continuous approximation

↳ We use the **Hard Concrete Distribution**: (relaxed version of discrete 0/1 values)

1. Instead of directly pruning weights (0 or 1), we assign a probability that a weight is active.
2. We sample values from a stretched sigmoid function ($\in [0, 1]$) to approximate binary decisions.
3. This allows gradient-based learning, making L_0 Regularization trainable.

• **Probabilistic Relaxation**: a mathematical trick used to approximate non-differentiable discrete decisions (e.g., selecting whether a weight should be active or zero) with continuous and differentiable approximations.

• **Why Concrete in the Hard Concrete Distribution?**

• The Concrete Distribution (or Relaxed Bernoulli Distribution) is a continuous approximation of a discrete Bernoulli (0/1) Distribution.

• **Why Hard in the Hard Concrete Distribution?** (most output nearly binary (0/1))

• Even though the function is continuous, the output strongly resembles a binary decision
⇒ The relaxation still maintains a "Hard" selection because most values get pushed close to 0 or 1.

The Hard Concrete Distribution modifies the Concrete Distribution by adding additional stretching or clipping.

Mathematical Trick: Hard Concrete Distribution

$$z = \sigma \left(\frac{\log(u) + \log(1-u) + \log(\alpha)}{\beta} \right) (\frac{\beta}{2} - \gamma) + \gamma$$

- $u \sim U(0,1)$, a sample from a Uniform $(0,1)$ Distribution

→ introduces stochasticity, allowing the model to explore different sparsity patterns

→ this randomness ensures that weight selection is a smooth process rather than an abrupt, discrete change

- $\log(u) + \log(1-u)$, (Logistic Noise), a transformation of u that maps it to a range from $-\infty$ to ∞

→ transforms Uniform D into a Logistic D , which helps in approximating binary choices in a smooth way

→ ensures that small differences in u leads to smooth changes in the final output z , making the function differentiable

- $\log(\alpha)$, a trainable parameter controlling how likely a weight is to be zeroed out (sparsity)

→ allows the model to learn which connections to remove dynamically during training

→ a large $\alpha \Rightarrow$ the weight is more likely to be pruned, a small $\alpha \Rightarrow$ the weight is more likely to stay

- $\beta \rightarrow$ Temperature Parameter (Softening the Decision), a scaling factor that controls the sharpness of a $0/1$ transition

→ controls how smooth or sharp the transition is from keeping a weight to pruning it

→ a large $\beta \Rightarrow$ smoother transitions (gradually learning sparsity), small $\beta \Rightarrow$ harder thresholding (close to 0)

- $\sigma(\cdot) \rightarrow$ Sigmoid Activation: $\sigma(x) = \frac{1}{1 + e^{-x}}$

→ smoothly maps the transformed value into the range $[0,1]$

→ ensures that the function is differentiable so gradient-based optimization works

- $(\frac{\beta}{2} - \gamma) + \gamma \rightarrow$ Stretching & Shifting, scales and shift the sigmoid output

→ ensures that the mask z can be in the range $[\frac{\beta}{2}, \frac{\beta}{2}]$ instead of just $[0,1]$

→ γ (negative shift): prevent weights from being too close to zero during training

→ $\frac{\beta}{2}$ (positive stretch): ensures that the function can reach exactly 1 when needed

Clamping z Returns 0 and 1: $z = \text{clamp}(z, 0, 1)$

→ a final step to force z to stay between 0 and 1

→ even with smooth transformations, z could sometimes exceed these limits, so we clip it

→ ensures z behaves like a proper probability mask

- Mask: a matrix M with values typically between $[0,1]$ or $\{0,1\}$, applied element-wise to another tensor X : $X' = M \odot X$

- $M_{ij} = 1$: the corresponding element X_{ij} is kept

- $M_{ij} = 0$: the corresponding element X_{ij} is zeroed out (ignored)

- $0 < M_{ij} < 1$: represents a probabilistic mask allowing gradient-based learning