

## Focus on $L_k$ Regularization

$$x = (x_1, \dots, x_n), \text{ For } q \in \mathbb{N}^*, \|x\|_q = \sqrt[q]{\sum |x_i|^q}$$
$$\begin{cases} q=0, \|x\|_0 = \sum \mathbb{1}_{x_i \neq 0} \end{cases}$$

INDICATOR FUNCTION:  $\mathbb{1}_{x_i} = \begin{cases} 0, & x_i = 0 \\ 1, & x_i \neq 0 \end{cases}$

## I - Focus on $L_0$ Regularization

$L_0$  Regularization refers to minimizing the number of non-zero weights in a model

↳ goal: remove unnecessary parameters and make the network sparse (sparse = interpretable & efficient)

↳ effect: **pruning** = forces some weights to be exactly zero

⇒ Because of its discontinuous nature: Standard Gradient Descent for Regularization UNUSABLE

$$X \quad \mathcal{L} = \underbrace{\text{Loss}(W)}_{\text{total loss}} + \underbrace{\lambda \|W\|_0}_{\substack{\text{regular loss from our NN (like MSE)}}} + \underbrace{\lambda \|W\|_0}_{L_0 \text{ penalty}}$$

⇒ We need a continuous approximation

↳ We use the **Hard Concrete Distribution**: (relaxed version of discrete 0/1 values)

1. Instead of directly pruning weights (0 or 1), we assign a probability that a weight is active.
2. We sample values from a stretched sigmoid function ( $\in [0, 1]$ ) to approximate binary decisions.
3. This allows gradient-based learning, making  $L_0$  Regularization trainable.

• **Probabilistic Relaxation**: a mathematical trick used to approximate non-differentiable discrete decisions (e.g., selecting whether a weight should be active or zero) with continuous and differentiable approximations.

• Why **Concrete** in the Hard Concrete Distribution?

• The Concrete Distribution (or Relaxed Bernoulli Distribution) is a continuous approximation of a discrete Bernoulli (0/1) Distribution.

• Why **Hard** in the Hard Concrete Distribution? (most output nearly binary 0/1)

• Even though the function is continuous, the output strongly resembles a binary decision  
⇒ The relaxation still maintains a "hard" selection because most values get pushed close to 0 or 1.

The Hard Concrete Distribution modifies the Concrete Distribution by adding additional stretching or clipping.

## Mathematical Trick: Hard Concrete Distribution

$$z = \frac{\log(u) - \log(1-u) + \log(\alpha)}{\beta} (\beta - \gamma) + \gamma$$

- $u \sim U(0,1)$ , a sample from a Uniform  $(0,1)$  Distribution
  - introduces stochasticity, allowing the model to explore different sparsity patterns
  - this randomness ensures that weight selection is a smooth process rather than an abrupt, discrete change
- $\log(u) - \log(1-u)$ , (Logistic Noise), a transformation of  $u$  that maps it to a range from  $-\infty$  to  $\infty$ 
  - Transforms Uniform  $D$  into a Logistic  $D$ , which helps in approximating binary choices in a smooth way
  - ensures that small differences in  $u$  leads to smooth changes in the final output  $z$ , making the function differentiable
- $\log(\alpha)$ , a trainable parameter controlling how likely a weight is to be zeroed out (sparsity)
  - allows the model to learn which connections to remove dynamically during training
  - a large  $\alpha \Rightarrow$  the weight is more likely to be pruned, a small  $\alpha \Rightarrow$  the weight is more likely to stay
- $\beta \rightarrow$  Temperature Parameter (softening the Decision), a scaling factor that controls the sharpness of a HP transition
  - controls how smooth or sharp the transition is from keeping a weight to pruning it
  - a large  $\beta \Rightarrow$  smoother transitions (gradually learning sparsity), small  $\beta \Rightarrow$  harder thresholding (close to 0)
- $\sigma(\cdot) \rightarrow$  Sigmoid Activation:  $\sigma(x) = \frac{1}{1+e^{-x}}$ 
  - smoothly maps the transformed value into the range  $[0,1]$
  - ensures that the function is differentiable so gradient-based optimization works
- $(\beta - \gamma) + \gamma \rightarrow$  stretching & shifting, scales and shift the sigmoid output
  - ensures that the mask  $z$  can be in the range  $[\gamma, \beta]$  instead of just  $[0,1]$
  - $\gamma$  (negative shift): prevent weights from being too close to zero during training
  - $\beta$  (positive stretch): ensures that the function can reach exactly 1 when needed

## Clamping $z$ Between 0 and 1: $z = \text{Clamp}(z, 0, 1)$

- a final step to force  $z$  to stay between 0 and 1
- even with smooth transformations,  $z$  could sometimes exceed these limits, so we clip it
- ensures  $z$  behaves like a proper probability mask

**Mask:** a matrix  $M$  with values typically between  $[0,1]$  or  $\{0,1\}$ , applied element-wise to another tensor  $X$ :  $X' = M \odot X$

- $M_{ij} = 1$ : the corresponding element  $X_{ij}$  is kept
- $M_{ij} = 0$ : the corresponding element  $X_{ij}$  is zeroed out (ignored)
- $0 < M_{ij} < 1$ : represents a probabilistic mask allowing gradient-based learning